
Machine Listening for Music and Sound Analysis

Lecture 2 – Machine Learning/Deep Learning

Dr.-Ing. Jakob Abeßer

Fraunhofer IDMT

jakob.abesser@idmt.fraunhofer.de

<https://www.machinelisting.de>

Learning Objectives

- Introduction
- Learning paradigms
- Machine learning (ML) project pipeline
- Deep learning

Introduction

■ Goals

- "...give computers the ability to learn without being explicitly programmed" [Samuels, 1959]
- Learning structures in given (un)labeled data to make predictions on new / unseen data

■ Paradigm change

- Before: manually designed / general-purpose features
- Now: joint representation learning (features) & data modeling (classification)


■ Related disciplines

- Statistics, data science, optimization

Introduction

Terminology

- Artificial Intelligence (AI)
 - “an agent’s ability to achieve goals in a wide range of environments” [Legg & Hutter, 2007]
- Machine Learning (ML)
 - Pattern recognition, data modeling, learning, prediction
- Deep Learning (DL)
 - (Brain-inspired) artificial neural networks (ANN)
- Data Science
 - Knowledge extraction from data



Introduction

Application Scenarios

- Computational finance (credit scoring, algorithmic trading)
- Computer vision (face & object recognition, motion detection)
- Computational biology (tumor detection, drug discovery, DNA sequencing)
- Energy (price & load forecasting)
- Predictive maintenance (automotive, aerospace, manufacturing)
- Natural language processing (sentiment classification, text search, translation)
- Machine listening (music transcription, instrument recognition, sound event detection, acoustic scene classification)

Learning Paradigms

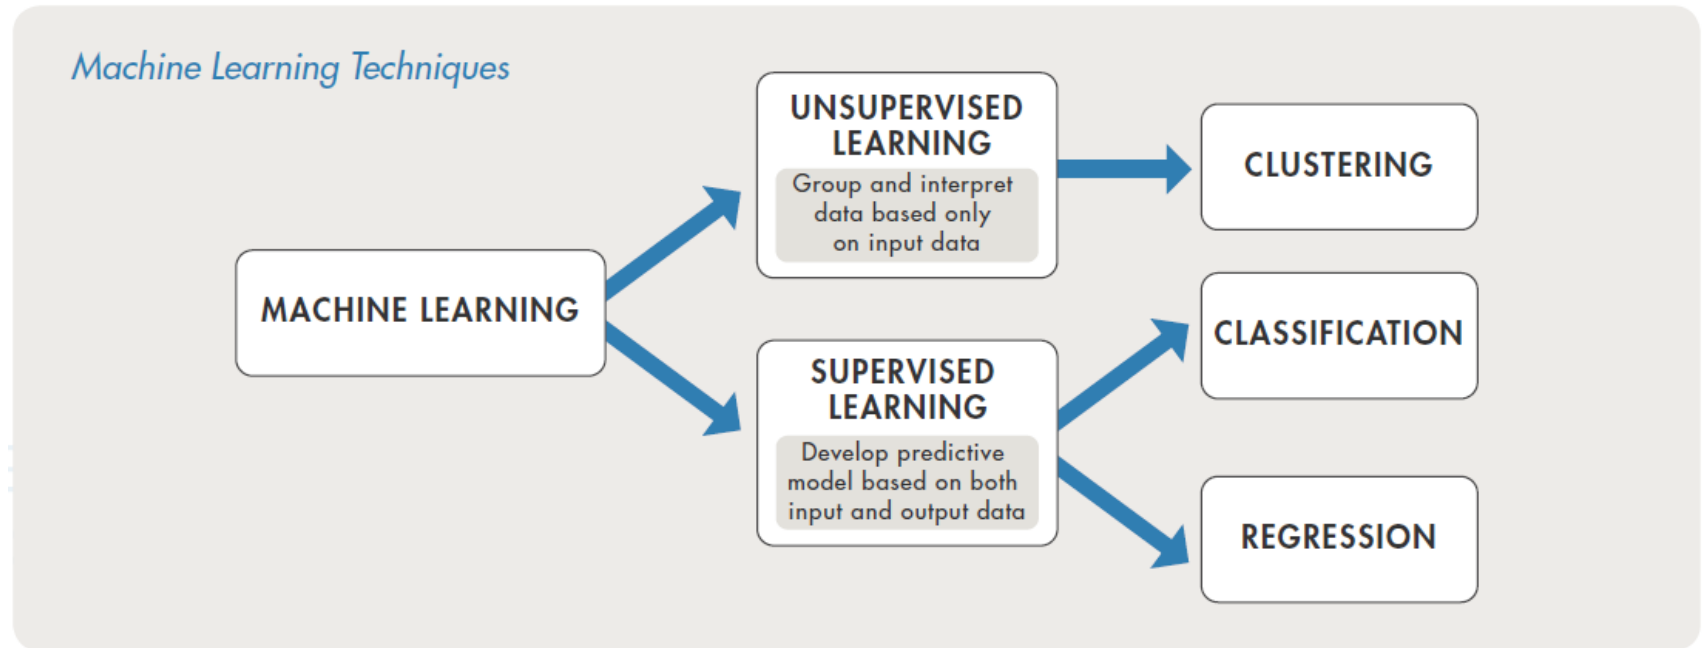


Fig. 1

Learning Paradigms

Unsupervised Learning

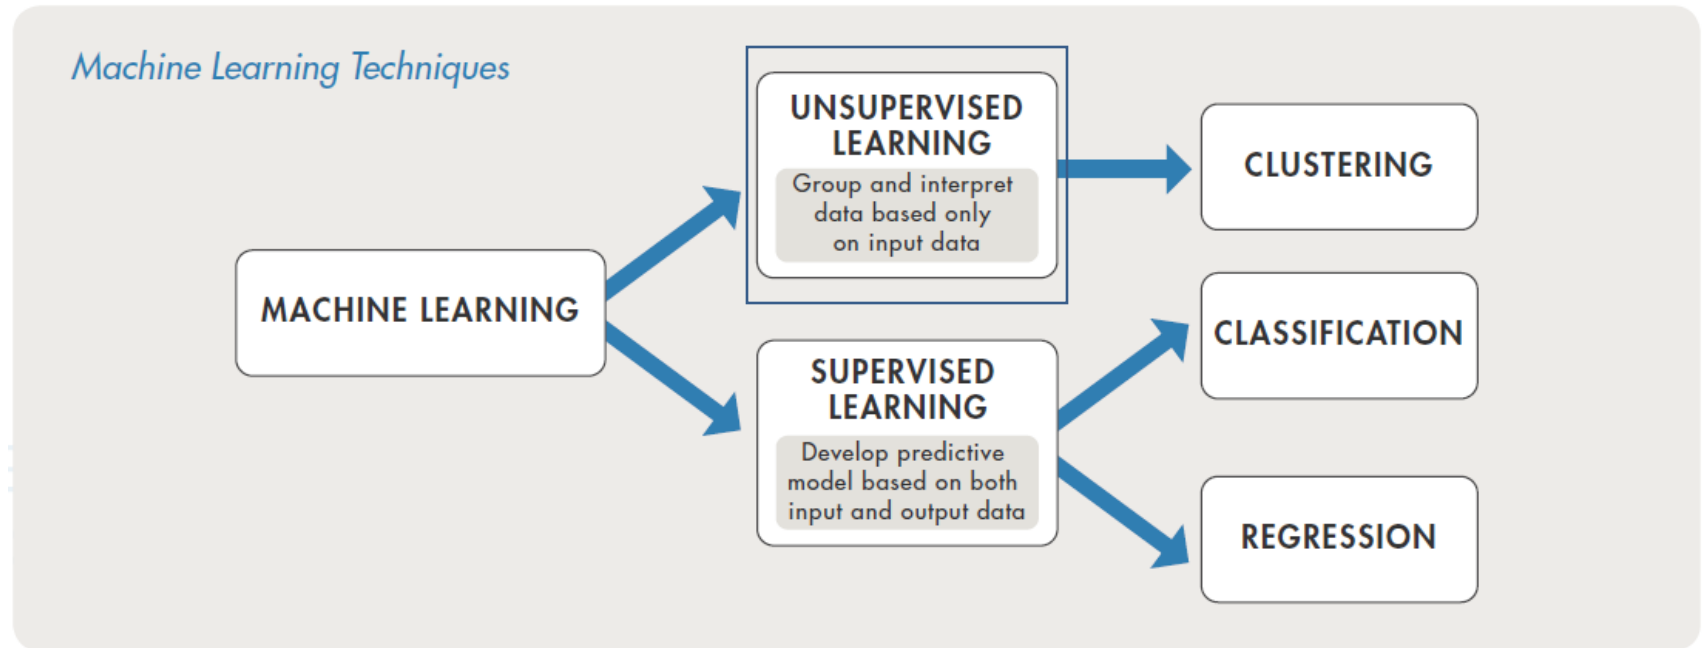


Fig. 1

Learning Paradigms


Unsupervised Learning

- Goal
 - Find hidden **structure** and **patterns** in data
 - **No annotations** available

Learning Paradigms

Unsupervised Learning


- Goal
 - Find hidden **structure** and **patterns** in data
 - **No annotations** available
- Clustering
 - **Grouping** of **similar** data instances



Learning Paradigms

Unsupervised Learning


- Goal
 - Find hidden **structure** and **patterns** in data
 - **No annotations** available
- Clustering
 - **Grouping** of **similar** data instances



Learning Paradigms

Unsupervised Learning

- Goal
 - Find hidden **structure** and **patterns** in data
 - **No annotations** available
- Clustering
 - **Grouping** of **similar** data instances




Learning Paradigms

Unsupervised Learning


■ Challenges

- What is the **optimal number of clusters**?



2 clusters

?




3 clusters

Learning Paradigms

Unsupervised Learning

- K -means clustering
 - Initialize K "means" randomly (=cluster centroids)




Learning Paradigms

Unsupervised Learning

■ K-means clustering

■ $K=3$




Learning Paradigms

Unsupervised Learning

- *K*-means clustering

- Assignment: assign each data point to its closest mean




Learning Paradigms

Unsupervised Learning

- *K*-means clustering

- Assignment: assign each data point to its closest mean




Learning Paradigms

Unsupervised Learning

- *K*-means clustering

- Assignment: assign each data point to its closest mean




Learning Paradigms

Unsupervised Learning

- K -means clustering

- Update: update mean by average over all assigned data points




Learning Paradigms

Unsupervised Learning

- K -means clustering

- Assignment: re-assign data points to closest mean




Learning Paradigms

Unsupervised Learning

- K -means clustering

- Update: re-assign data points to closest mean (repeat until convergence)



Learning Paradigms

Supervised Learning

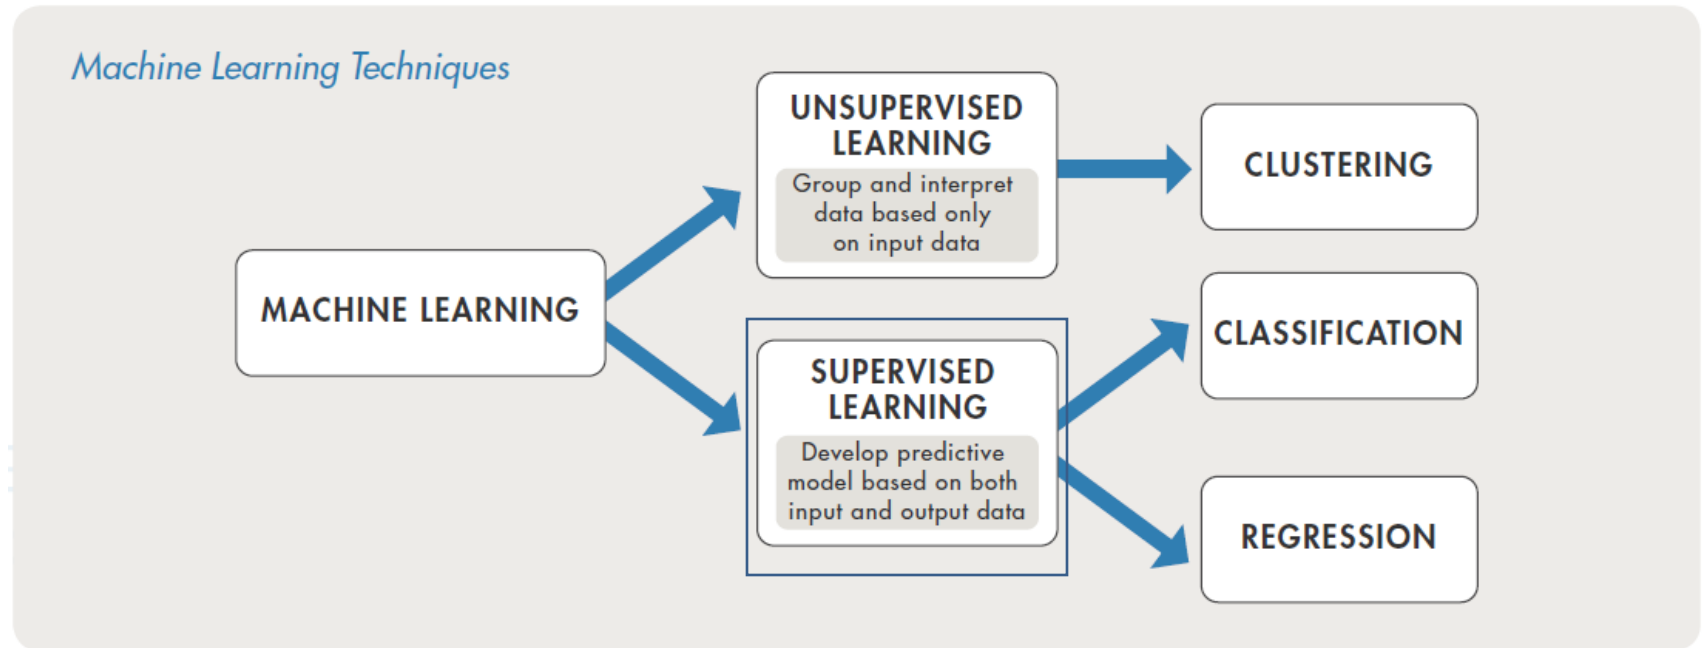
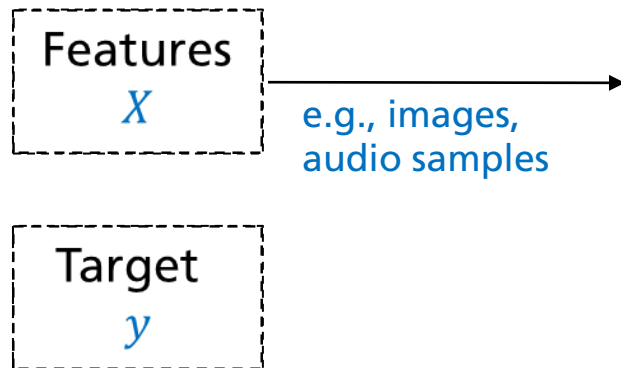


Fig. 1

Learning Paradigms

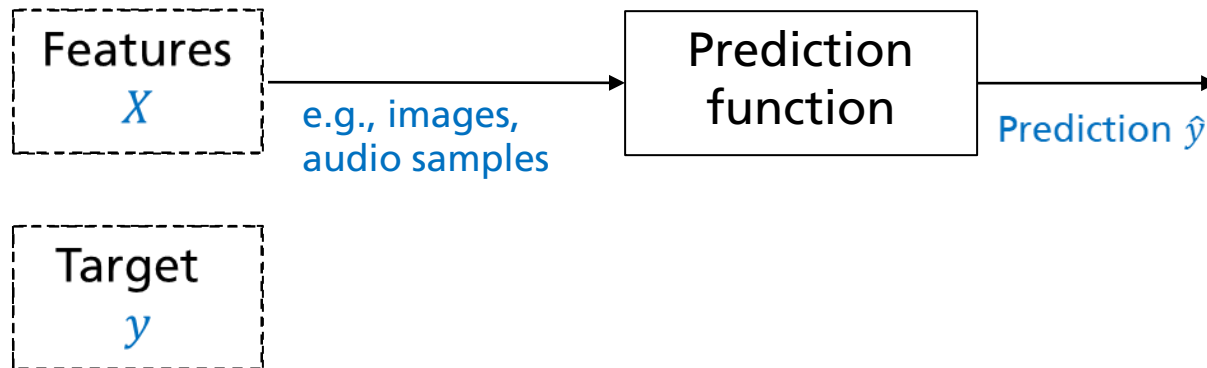
Supervised Learning



Own

Learning Paradigms

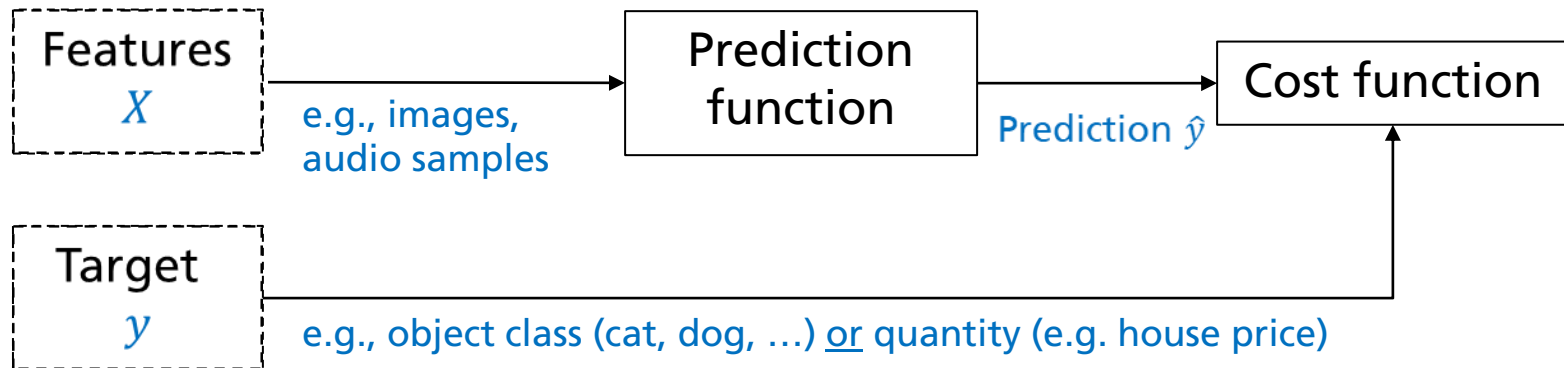
Supervised Learning



Own

Learning Paradigms

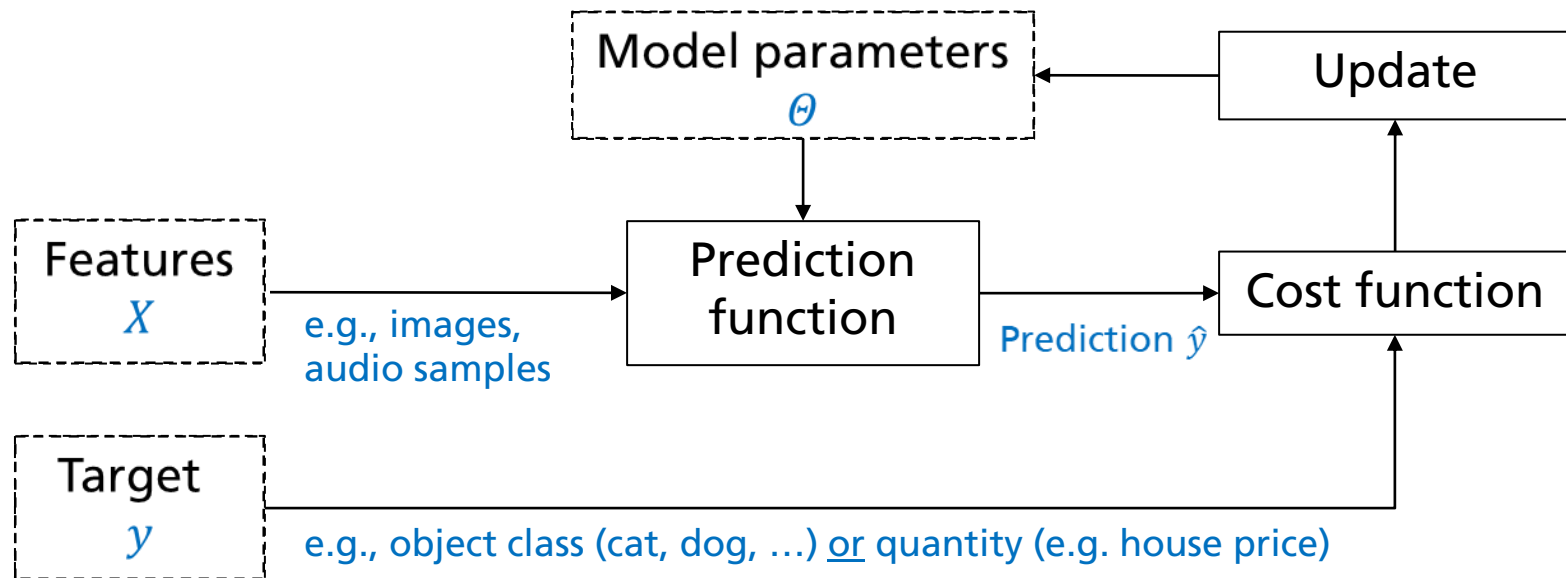
Supervised Learning



Own

Learning Paradigms

Supervised Learning



Own

Learning Paradigms

Supervised Learning - Classification

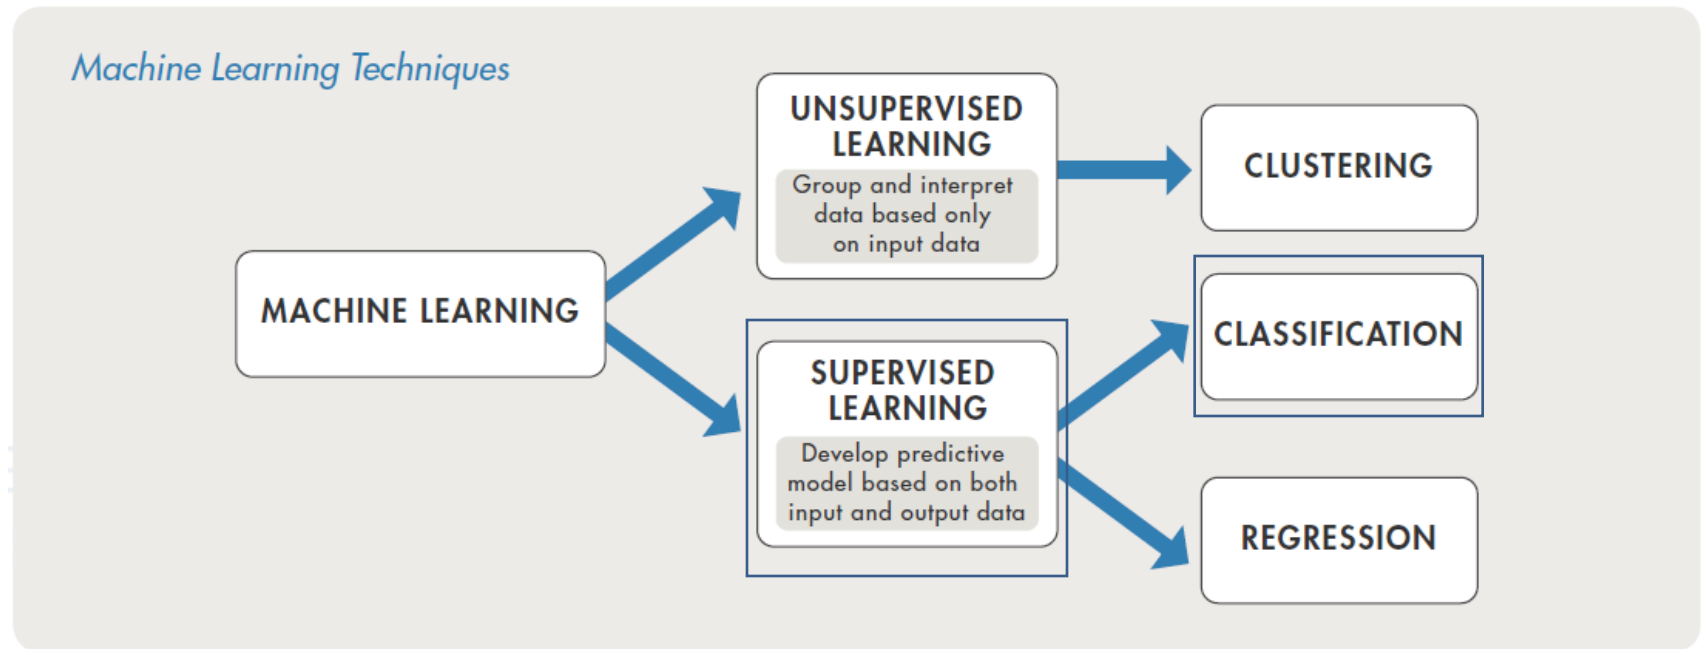


Fig. 1

Learning Paradigms

Supervised Learning - Classification

- Predict one or multiple categorical labels from features
 - Examples → music genre, instrument(s), key

Learning Paradigms

Supervised Learning - Classification

- Predict one or multiple categorical labels from features
 - Examples → music genre, instrument(s), key
- Feature space modeling (Example: 2 classes)

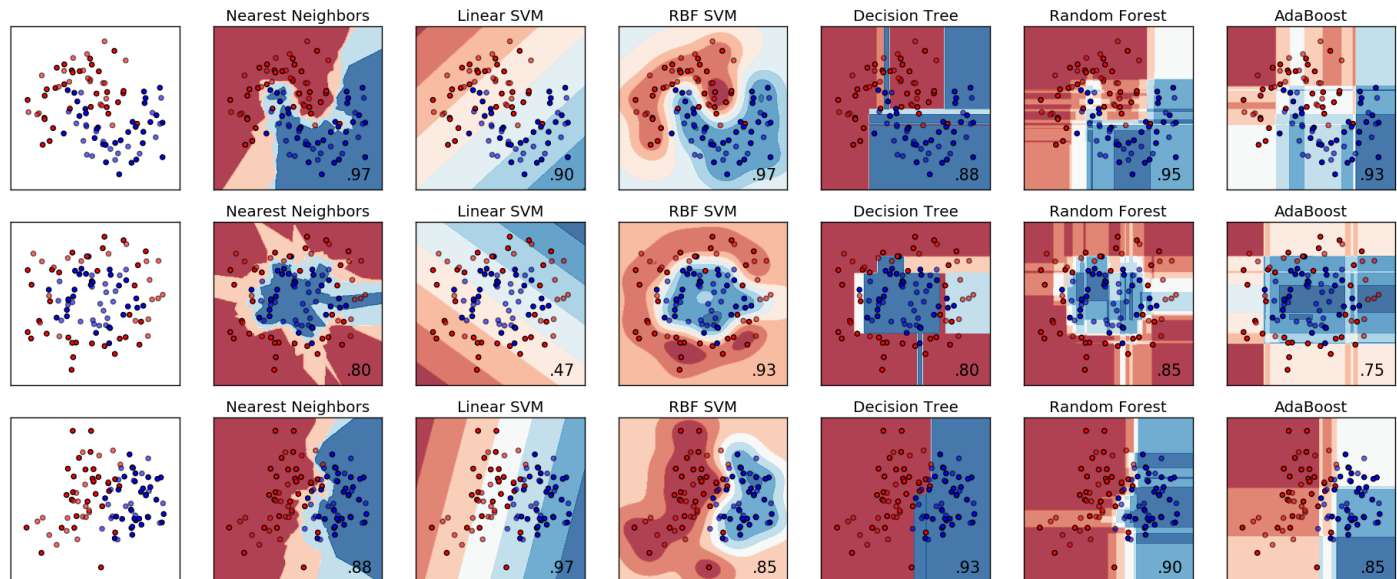


Fig. 3

Learning Paradigms

Supervised Learning - Classification

- Example: k -Nearest Neighbors
 - Training → Store all examples

Learning Paradigms

Supervised Learning - Classification

- Example: k -Nearest Neighbors
 - Training → Store all examples





Fig. 4

Learning Paradigms

Supervised Learning - Classification


- Example: k -Nearest Neighbors
 - Training → Store all examples
 - Test → Assign test item to dominant class label of the k closest training data items



Learning Paradigms

Supervised Learning - Classification

- Example: k -Nearest Neighbors
 - Training → Store all examples
 - Test → Assign test item to dominant class label of the k closest training data items



Learning Paradigms

Supervised Learning - Classification

- Example: k -Nearest Neighbors
 - Training → Store all examples
 - Test → Assign test item to dominant class label of the k closest training data items
- Distance measures
 - Euclidean distance, Manhattan distance, cosine distance, ...

Learning Paradigms

Supervised Learning

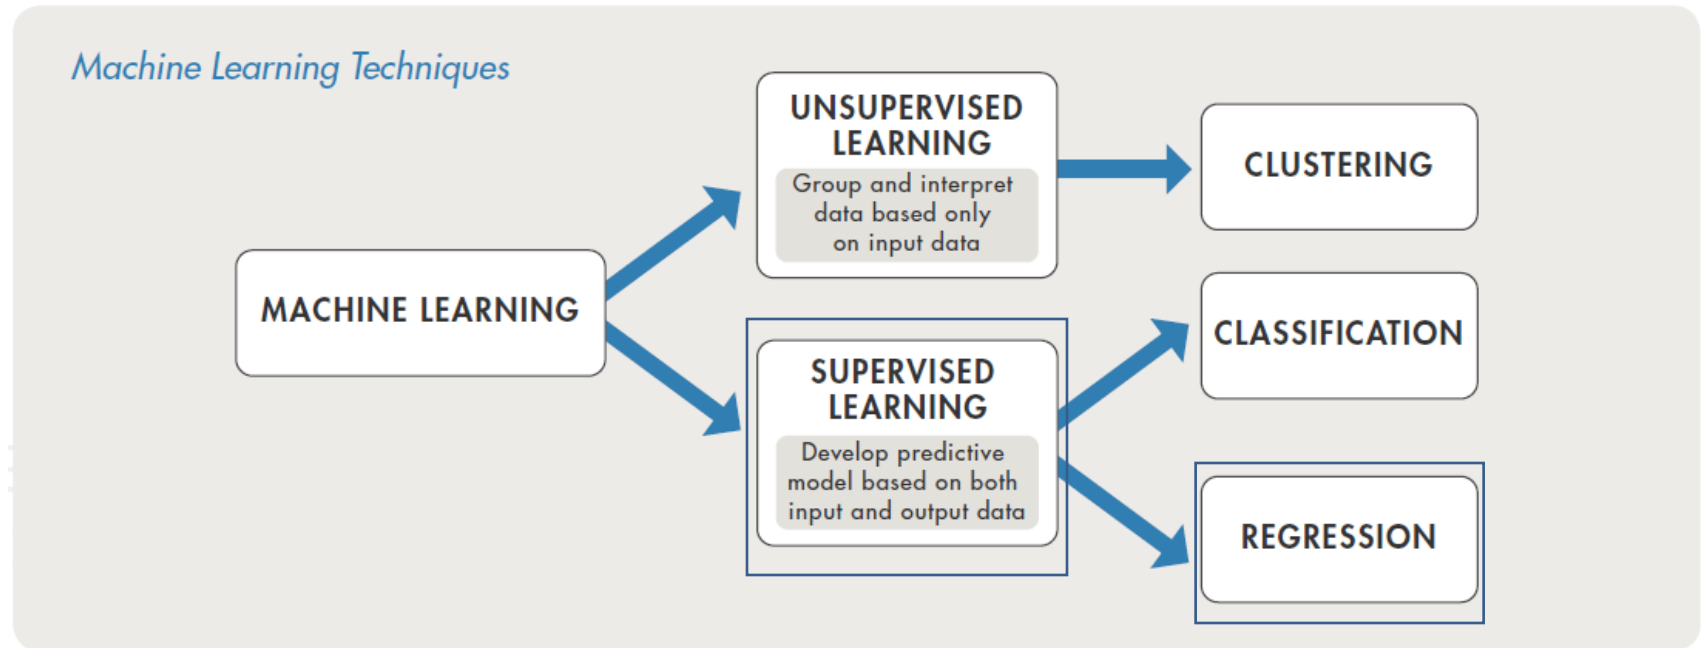


Fig. 1

Learning Paradigms

Supervised Learning - Regression

■ Goal

- Predict a dependent (response) variable given one or multiple independent variables (features)
- Continuous quantities

■ Examples

- Univariate (linear) regression:
 - $y \approx \beta_0 + \beta_1 x_1$
 - $\beta_0 \rightarrow$ bias
 - $\beta_1 \rightarrow$ weight

Learning Paradigms

Supervised Learning - Regression


■ Goal

- Predict a dependent (response) variable given one or multiple independent variables (features)
- Continuous quantities

■ Examples

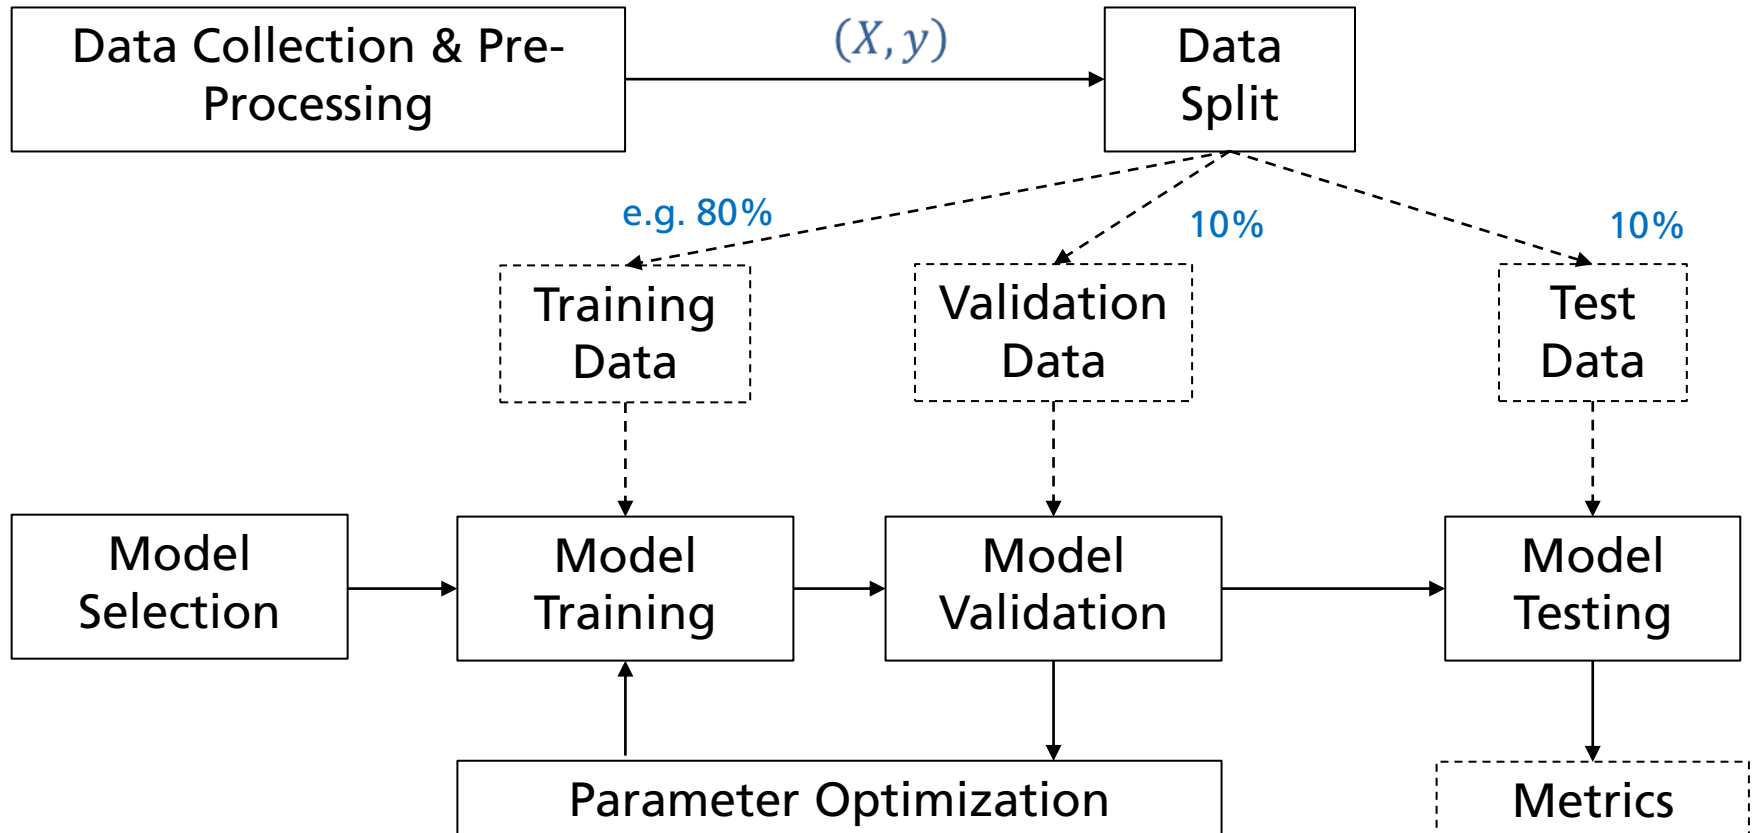
■ Univariate (linear) regression:

- $y \approx \beta_0 + \beta_1 x_1$
 - $\beta_0 \rightarrow$ bias
 - $\beta_1 \rightarrow$ weight



Own

ML Project Pipeline Overview



ML Project Pipeline

Data Split

- Training Set ■
- Model learns from this data



ML Project Pipeline

Data Split

- Training Set ■
 - Model learns from this data
- Validation / Development Set ■
 - Used to fine-tune the model (hyper)parameters
 - Model occasionally sees but does not learn from this data



ML Project Pipeline




Data Split

- Training Set ■
 - Model learns from this data
- Validation / Development Set ■
 - Used to fine-tune the model (hyper)parameters
 - Model occasionally sees but does not learn from this data
- Test set ■
 - Only used once after the model training & tuning is completed
 - Should reflect the targeted real-world use case for the model



ML Project Pipeline

Data Split

- Training Set 
 - Model learns from this data
- Validation / Development Set 
 - Used to fine-tune the model (hyper)parameters
 - Model occasionally sees but does not learn from this data
- Test set 
 - Only used once after the model training & tuning is completed
 - Should reflect the targeted real-world use case for the model
- Common split ratios
 - 80/10/10% or even 98/1/1% (for large datasets)



ML Project Pipeline

Data Collection & Pre-Processing

- Data collection
 - Check for available data resources for given (or related) task
 - Collect / record / annotate new data
 - Ensure data variability
 - Example (from acoustic condition monitoring) → include different motor engine types & conditions, recording locations, microphones, ...

ML Project Pipeline

Data Collection & Pre-Processing

- Data collection
 - Check for available data resources for given (or related) task
 - Collect / record / annotate new data
 - Ensure data variability
 - Example (from acoustic condition monitoring) → include different motor engine types & conditions, recording locations, microphones, ...
- Data cleanup / pre-processing
 - Remove errors, silence, empty files, ...
 - Balance dataset (proportions among class examples)
 - Normalize (depends on the model)

ML Project Pipeline

Model Selection

- Many models and approaches exist
 - Types (SVM, GMM, logistic regression, DNNs)
 - Hyperparameters (SVM kernel functions, DNN layer types)

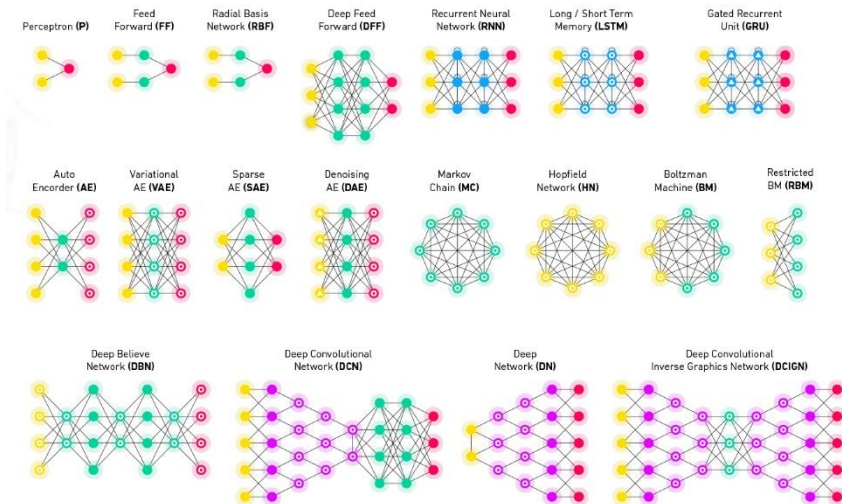


Fig. 6

ML Project Pipeline

Model Selection

- Many models and approaches exist
 - Types (SVM, GMM, logistic regression, DNNs)
 - Hyperparameters (SVM kernel functions, DNN layer types)
- Often constrained by the use-case / task
 - Model complexity (memory, training time, training data amount)




Fig. 6

ML Project Pipeline

Model Selection

- Many models and approaches exist
 - Types (SVM, GMM, logistic regression, DNNs)
 - Hyperparameters (SVM kernel functions, DNN layer types)
- Often constrained by the use-case / task
 - Model complexity (memory, training time, training data amount)
- Feature pre-processing depends on model type
- Use simple models for simple tasks

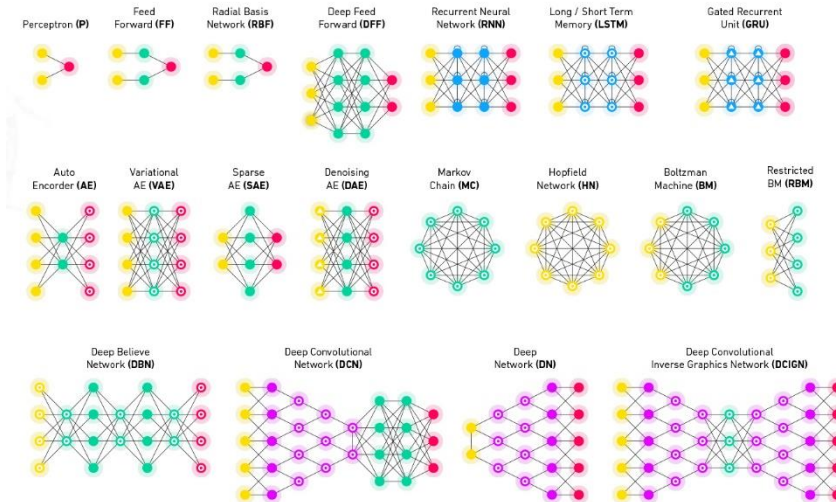


Fig. 6

ML Project Pipeline

Model Training

- Iterative process
 - Typically: start with random parameter initialization

ML Project Pipeline

Model Training

- Iterative process
 - Typically: start with random parameter initialization
 - Use (batches of) training data to iteratively improve model predictions (optimization)
 - Learn from examples

ML Project Pipeline

Model Training

- Iterative process
 - Typically: start with random parameter initialization
 - Use (batches of) training data to iteratively improve model predictions (optimization)
 - Learn from examples
 - Update model parameters according to loss function

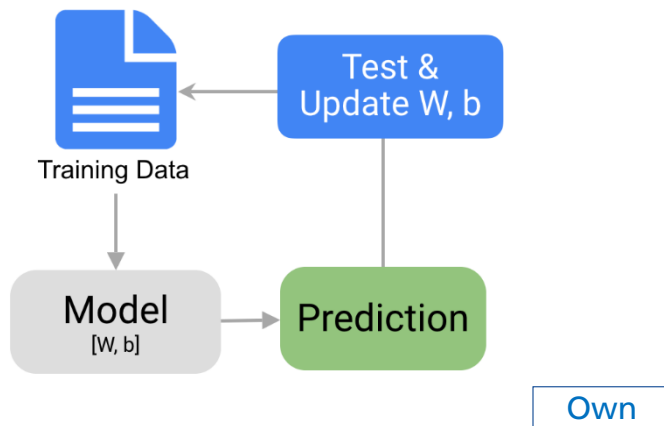
ML Project Pipeline

Model Training

- Example: linear regression

$$y \approx \beta_0 + \beta_1 x_1$$

- Training loop



ML Project Pipeline

Model Validation

- Regular model evaluation each or multiple training iteration

ML Project Pipeline

Model Validation

- Regular model evaluation each or multiple training iteration






Fig. 7

ML Project Pipeline

Model Validation

- Regular model evaluation each or multiple training iteration
- Helps to
 - optimize model (hyper)parameters
 - detect overfitting on training data
 - stop the training



Own

ML Project Pipeline

Model Testing

- Example: Binary classification evaluation

- True/false positives (TP/FP)
- True/false negatives (TN/FN)

		Prediction	
		1	0
Annotation	1	TP <i>true positives</i>	FN <i>false negatives</i>
	0	FP <i>false positives</i>	TN <i>true negatives</i>

Fig. 8

ML Project Pipeline

Model Testing

■ Example: Binary classification evaluation

■ True/false positives (TP/FP)

■ True/false negatives (TN/FN)

■ Metrics

■ Precision

■ Recall

■ Accuracy

■ F-score


		Prediction		
		1	0	
Annotation	1	TP <i>true positives</i>	FN <i>false negatives</i>	True Positive Rate Sensitivity Recall $R = \frac{TP}{TP+FN}$
	0	FP <i>false positives</i>	TN <i>true negatives</i>	False Positive Rate $FPR = \frac{FP}{FP+FN}$ Specificity $\text{Specificity} = \frac{TN}{FP+FN}$
		Precision $P = \frac{TP}{TP+FP}$		Accuracy $ACC = \frac{TP+TN}{TP+TN+FP+FN}$

Fig. 8

Deep Learning

Introduction

- Artificial neural networks → mimic brain processing
 - Connected neurons
 - Weighted input summation
 - Non-linear processing



Deep Learning

Introduction

- Artificial neural networks → mimic brain processing
 - Connected neurons
 - Weighted input summation
 - Non-linear processing
- Shallow networks

Simple Neural Network





Fig. 9


Fig. 10

Deep Learning

Introduction


- Artificial neural networks → mimic brain processing
 - Connected neurons
 - Weighted input summation
 - Non-linear processing
- Shallow networks → deep networks

Simple Neural Network



● Input Layer

Deep Learning Neural Network



● Hidden Layer

● Output Layer





Fig. 9

Fig. 10

Deep Learning

Introduction

- Hierarchical feature learning
 - Example (face recognition)



Edges, curves

Fig. 11


First layers

Final layers


Deep Learning

Introduction

- Hierarchical feature learning
 - Example (face recognition)



Edges, curves



Shapes, object parts

Fig. 11


First layers

Final layers


Deep Learning

Introduction

- Hierarchical feature learning
 - Example (face recognition)



Edges, curves



Shapes, object parts



Objects (faces)

Fig. 11


First layers

Final layers

Deep Learning

Fully-connected (Deep) Neural Networks

Input layer




Own

Deep Learning

Fully-connected (Deep) Neural Networks

Input layer

Hidden layers



- 1) Weights w & biases b
- 2) Non-linear activation function $g(z)$

W_1, b_1

$$z_1 = W_1 x + b_1$$
$$a_1 = g(z_1)$$


Own

Deep Learning

Fully-connected (Deep) Neural Networks

Input layer

Hidden layers



1) Weights w & biases b
2) Non-linear activation function $g(z)$

$$\begin{aligned} z_1 &= W_1 x + b_1 & z_2 &= W_2 a_1 + b_2 \\ a_1 &= g(z_1) & a_2 &= g(z_2) \end{aligned}$$


Own

Deep Learning

Fully-connected (Deep) Neural Networks

Input layer


Hidden layers



Own

Deep Learning

Fully-connected (Deep) Neural Networks




Deep Learning


Activation Functions

- Activation functions add non-linearity
- Make networks more powerful in (complex) pattern recognition
- Examples:

Hyper Tangent Function



ReLU Function



Sigmoid Function

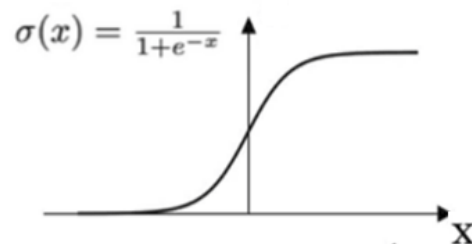


Fig. 12

Deep Learning Training

■ Overview

Features

X


Target

y

Own

Deep Learning Training


■ Overview



Own

Deep Learning Training

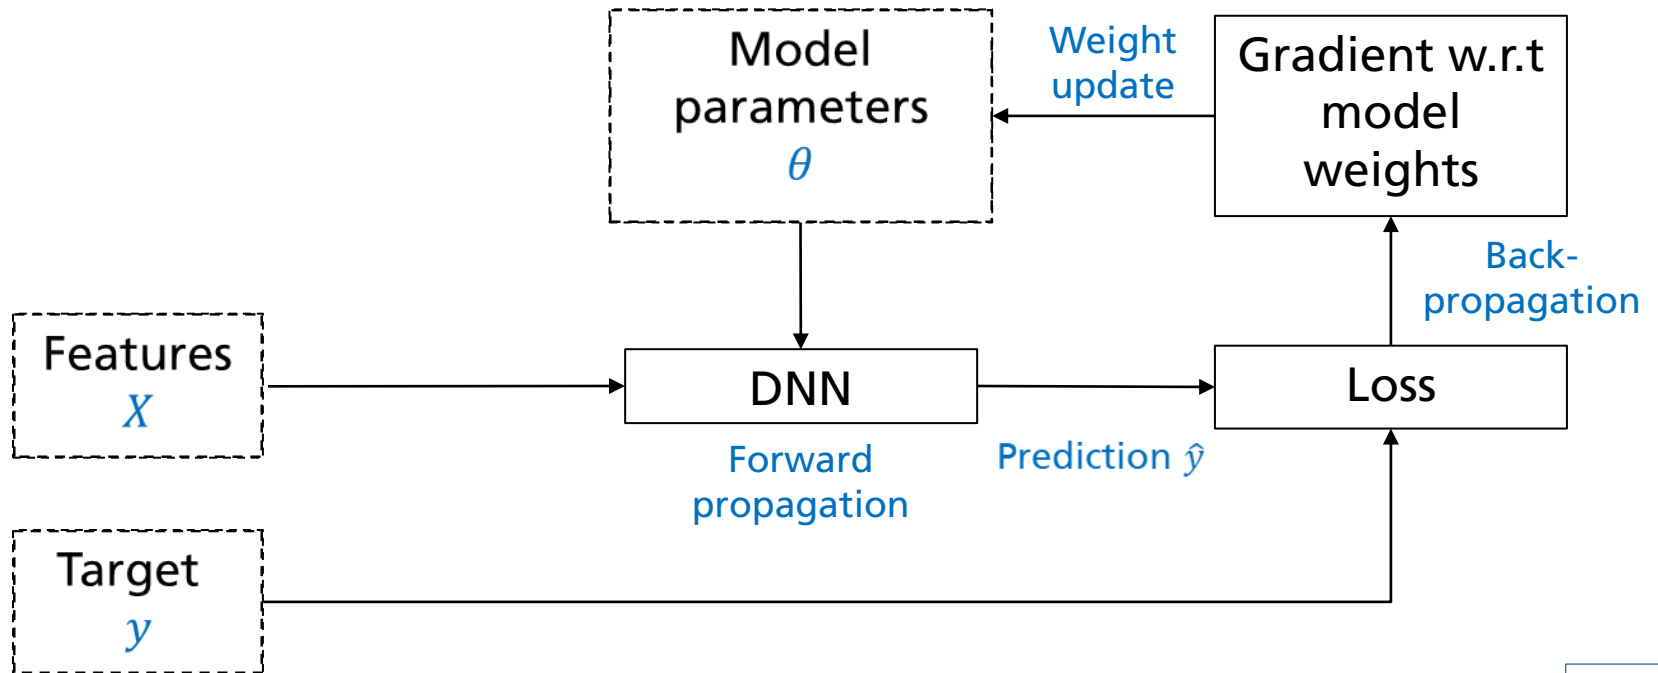
■ Overview



Own

Deep Learning Training


■ Overview



Own

Deep Learning Training

- Forward propagation → propagate batch of training data through the network → compute loss (compare to targets)



$$L(z, y) = - \left[y \log(z) + (1 - y) \log(1 - z) \right]$$

Fig. 20

Deep Learning Training

- Forward propagation → propagate batch of training data through the network → compute loss (compare to targets)
- Backpropagation → backpropagate loss → compute gradients of loss w.r.t. weights




Fig. 20

Deep Learning Training

- Forward propagation → propagate batch of training data through the network → compute loss (compare to targets)
- Backpropagation → backpropagate loss → compute gradients of loss w.r.t. weights
- Weights update → use gradients & learning rate to update weights


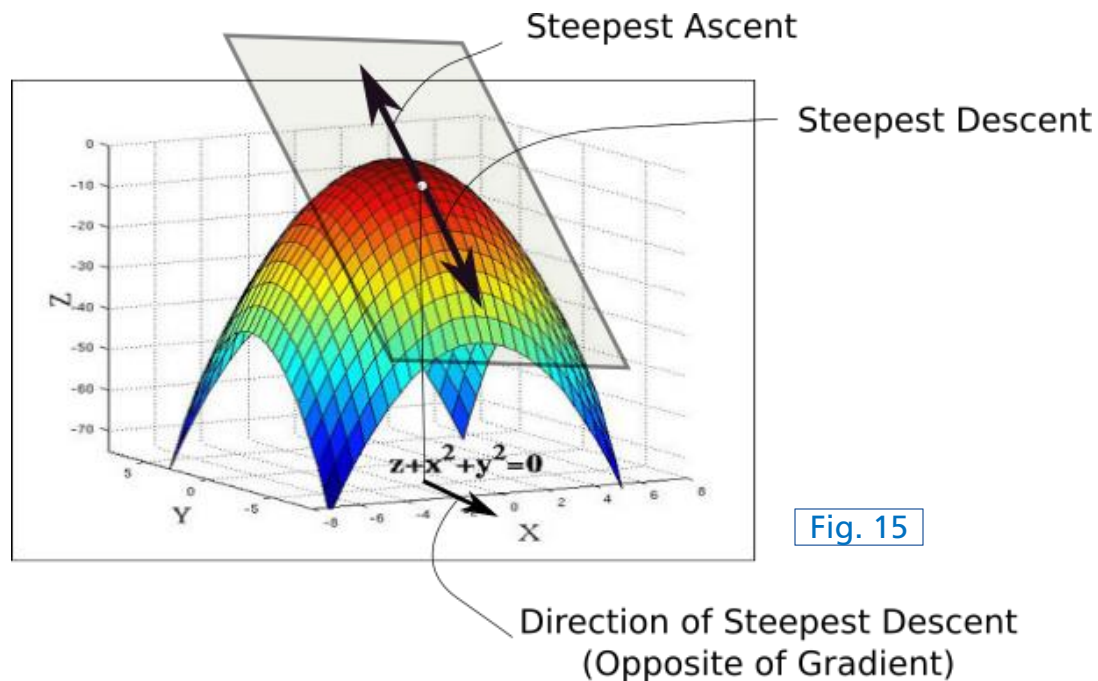


Fig. 20

Deep Learning Training


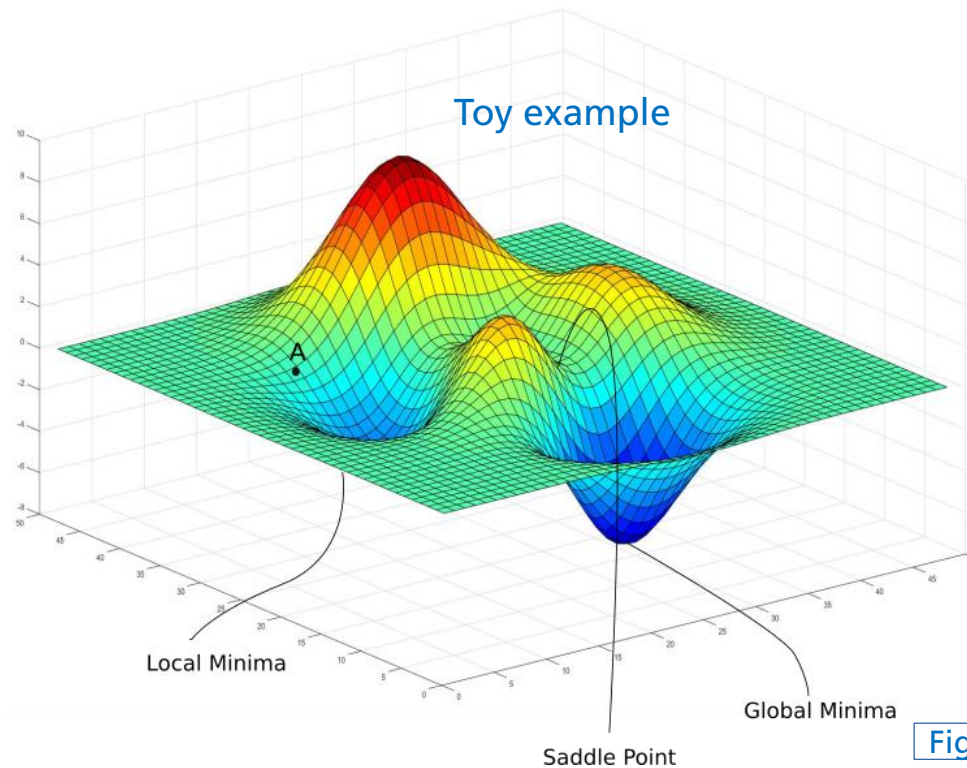
- Gradient descent
 - Move in opposite direction of gradient
 - Learning rate effects step size



Deep Learning Training

- Loss contour

- Goal → find global minima



Deep Learning Playground

- A neural network playground!
 - <https://playground.tensorflow.org>

Deep Learning

Convolutional Neural Networks (CNN)

- Convolutional layers

- "Convolution" → (local) dot-product between filter and input




Fig. 16

Deep Learning

Convolutional Neural Networks (CNN)

- Convolutional layers
 - “Convolution” → (local) dot-product between filter and input
 - Shared weights (across input)




Fig. 16

Deep Learning

Convolutional Neural Networks (CNN)

- Convolutional layers

- "Convolution" → (local) dot-product between filter and input
- Shared weights (across input)
- translation of input → translation of activations (equivariance)




Fig. 16

Deep Learning

Convolutional Neural Networks (CNN)

■ Convolutional layers

- "Convolution" → (local) dot-product between filter and input
- Shared weights (across input)
- translation of input → translation of activations (equivariance)

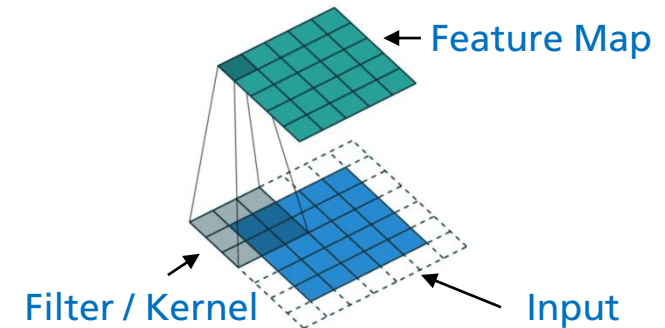


Fig. 16

■ Pooling → local aggregation / down-sampling

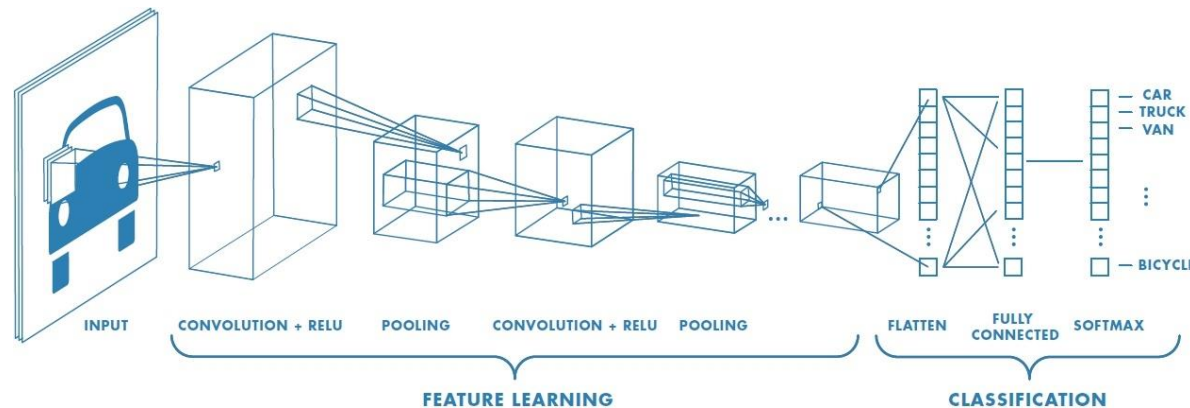


Fig. 17

Deep Learning

Recurrent Neural Networks (RNN)

- Recurrent layers
 - Model sequential data → model dynamic temporal behaviour
 - Internal memory state(s) → memorize previous data for future predictions

Deep Learning

Recurrent Neural Networks (RNN)

- Recurrent layers

- Model sequential data → model dynamic temporal behaviour
- Internal memory state(s) → memorize previous data for future predictions

- Vanishing gradient problem

- Gating mechanisms (Gated Recurrent Units (GRU), Long Short-term Memory (LSTM))




Fig. 18

Deep Learning

Recurrent Neural Networks (RNN)

- Application Examples

- One-to-many: sequential music generation (given a starting note)




Fig. 19

Deep Learning

Recurrent Neural Networks (RNN)

■ Application Examples

- One-to-many: sequential music generation (given a starting note)
- Many-to-one: sentiment classification (positive vs. negative)




Fig. 19

Deep Learning

Recurrent Neural Networks (RNN)

■ Application Examples

- One-to-many: sequential music generation (given a starting note)
- Many-to-one: sentiment classification (positive vs. negative)
- Many-to-many: machine translation (e.g., Spanish to German)





Fig. 19

Deep Learning

Autoencoders

- Symmetric architecture (decoder & encoder)




Own

Deep Learning

Autoencoders


- Symmetric architecture (decoder & encoder)
- Objective: minimize reconstruction error (e.g., mean squared error, MSE)



Deep Learning

Autoencoders

- Symmetric architecture (decoder & encoder)
- Objective: minimize reconstruction error (e.g., mean squared error, MSE)
- Compression of input (embedding)
- Prioritize important information → learn useful representations



Summary

- Introduction
 - Terminology, application scenarios
- Learning Paradigms
 - Unsupervised, supervised, self-supervised learning
- ML project pipeline
 - Data collection, pre-processing, split
 - Model selection, training, validation, testing
- Deep Learning
 - DNN, CNN, RNN, Autoencoders

References

- Introducing Machine Learning*. (2016). Retrieved from https://www.mathworks.com/content/dam/mathworks/tag-team/Objects/i/88174_92991v00_machine_learning_section1_ebook.pdf
- S. Legg, M. Hutter (2007). Universal Intelligence: A Definition of Machine Intelligence. *Minds & Machines*. 17 (4): 391-444.
- L. Samuel (1959). Some studies in machine learning using the game of checkers. *IBM Journal of research and development*. 3(3), 210-229
- Srihari, S. N. (2020). *Forward Propagation and Backward Propagation (Deep Learning Lecture)*. Retrieved from [https://cedar.buffalo.edu/~srihari/CSE676/6.5.0 Forward Backward.pdf](https://cedar.buffalo.edu/~srihari/CSE676/6.5.0%20Forward%20Backward.pdf)
- Virtanen, T., Plumbley, M. D., & Ellis, D. (Eds.). (2018). *Computational Analysis of Sound Scenes and Events*. Cham, Switzerland: Springer International Publishing.

Images

Fig. 1: [Machine Learning, 2016], p. 4, Fig. 2

Fig. 2: <https://i0.wp.com/www.sthda.com/sthda/RDoc/figure/clustering/partitioning-cluster-analysis-k-means-plot-4-groups-1.png>

Fig. 3: <https://i.stack.imgur.com/hsilO.png> (https://scikit-learn.org/stable/auto_examples/classification/plot_classifier_comparison.html)

Fig. 4: https://miro.medium.com/max/975/1*OyYyr9qY-w8RkaRh2TKo0w.png (reproduced)

Fig. 5: <https://lilianweng.github.io/lil-log/assets/images/self-sup-lecun.png>

Fig. 6: <https://www.asimovinstitute.org/wp-content/uploads/2019/04/NeuralNetworkZoo20042019.png>

Fig. 7: <https://www.educative.io/api/edpresso/shot/6668977167138816/image/5033807687188480>

Fig. 8: [Virtanen, 2018], p. 170, Fig. 6.7

Fig. 9: https://miro.medium.com/max/915/1*SJPacPhP4KDEB1AdhOFy_Q.png

Fig. 10: https://www.skampakis.com/wp-content/uploads/2018/03/simple_neural_network_vs_deep_learning.jpg

Fig. 11: https://pic4.zhimg.com/80/v2-057b248288a8af2f01272a956f862873_1440w.png

Fig. 12: https://blog.e-kursy.it/deeplearning4j-workshop/video/html/presentation_specific/img/4_activation_functions.png

Images

Fig. 13: <https://blog.paperspace.com/content/images/2018/05/challenges-1.png>

Fig. 14: <https://www.cs.umd.edu/~tomg/img/landscapes/noshort.png>

Fig. 15: <https://blog.paperspace.com/content/images/2018/05/grad.png>

Fig. 16: <https://www.wandb.com/articles/intro-to-cnns-with-wandb>

Fig. 17: <https://www.freecodecamp.org/news/an-intuitive-guide-to-convolutional-neural-networks-260c2de0a050/>

Fig. 18: <https://wiki.tum.de/download/attachments/22578349/RNN1.png>

Fig. 19: <https://stanford.edu/~shervine/teaching/cs-230/illustrations/architecture-rnn-ltr.png>

Fig. 20: [Srihari, 2020], p.8, (Fig. 1)

Thank you!

■ Any questions?

Dr.-Ing. Jakob Abeßer

Fraunhofer IDMT

Jakob.abesser@idmt.fraunhofer.de

<https://www.machinelistinging.de>