Machine Listening for Music and Sound Analysis

Lecture 4 – Music Information Retrieval II

Dr.-Ing. Jakob Abeßer

Fraunhofer IDMT

Jakob.abesser@idmt.fraunhofer.de

https://machinelistening.github.io

Fraunhofer IDMT

Overview

- Pitch Detection
- Instrument Recognition
- Source Separation

Pitch Detection
Introduction

- Pitch
  - (Subjective) psychoacoustic attribute of sound
  - Allows ordering from low to high in a frequency-related scale
    - Pitch ≠ frequency !
- Two subtasks

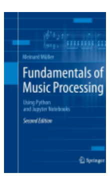

FMP Notebooks

1) Pitch detection

2) Voicing detection

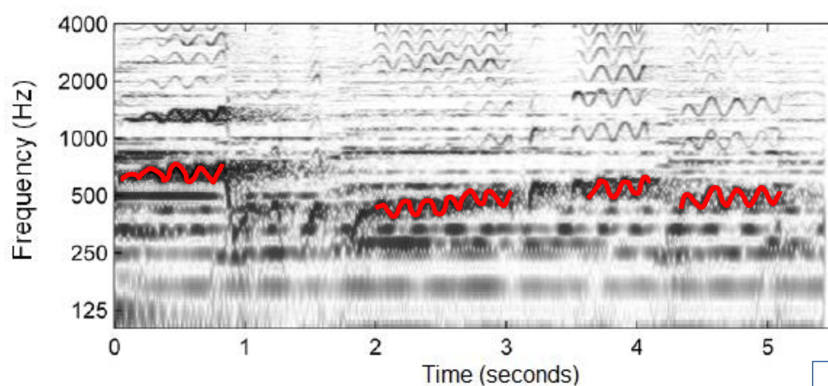

Fig. 1

Pitch Detection
Application Scenarios

- Music Instrument Tuning
- Music Education
- Music Transcription
- Bird Recognition

Fig. 2

Fig. 3

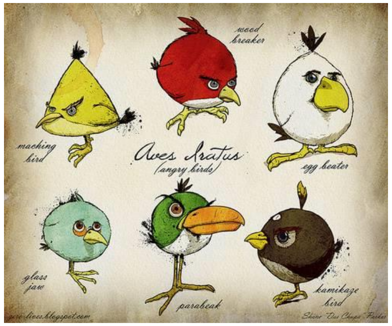

Fig. 4

Fraunhofer IDMT

Pitch Detection
Tasks

- Pitch detection of isolated monophonic instruments
- Predominant melody extraction in polyphonic music
- Polyphonic melody extraction

Increasing Difficulty & Complexity

Pitch Detection
Traditional Methods

- MELODIA [Salamon & Gomez, 2012]
  - Melody Extraction from polyphonic audio
- Steps
  - Sinusoid Extraction
    - Equal loudness filter
    - STFT
    - Detection of predominant peaks
    - Frequency refinement via instantaneous frequency (IF)

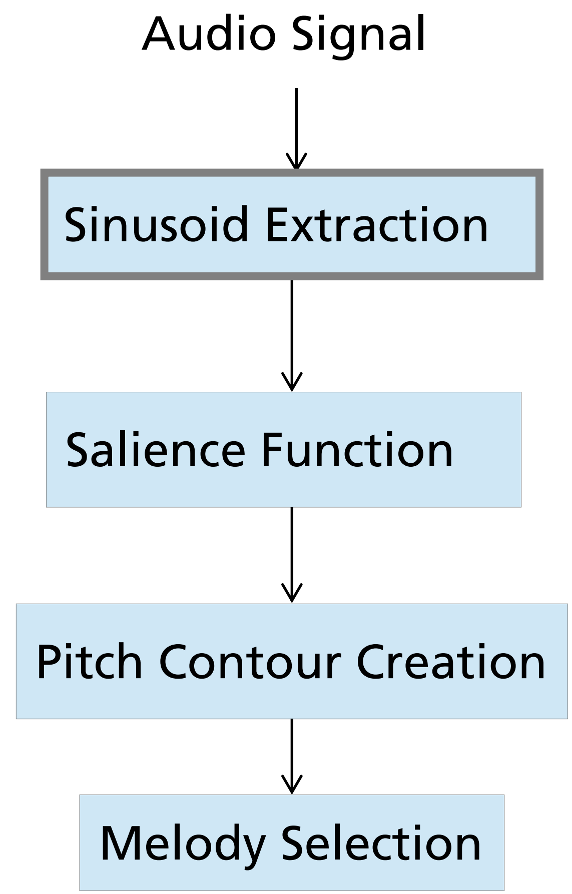

Pitch Detection
Traditional Methods

- Salience Function
    - Harmonic summation
        - Sum over possibile harmonic frequencies
        - Frequencies → pitch candidates

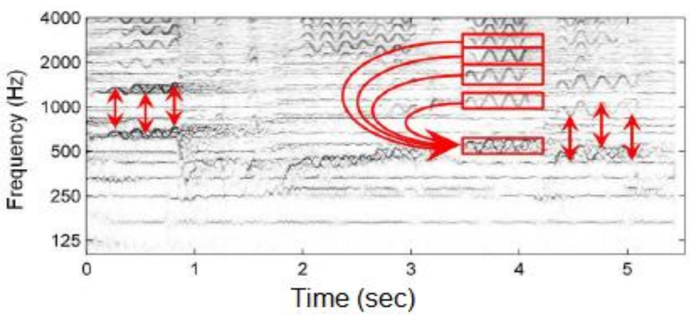

Fig. 5

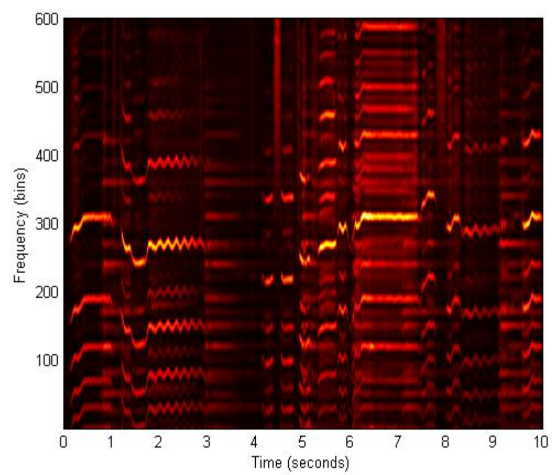

Audio Signal

↓

Sinusoid Extraction

↓

Salience Function

↓

Pitch Contour Creation

↓

Melody Selection

Pitch Detection
Traditional Methods

- Pitch contour creation & melody selection
    - Auditory streaming cues → group peaks to continuous paths (pitch contours)
    - Select melody contours using features (e.g. average pitch / salience, vibrato)
    - Note formation (one pitch value)

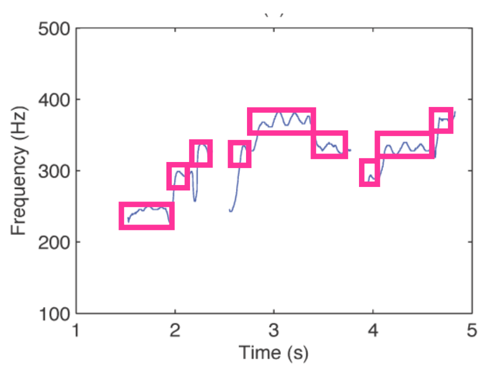

Pitch contour(s)

↓

Note events

Audio Signal

↓

Sinusoid Extraction

↓

Salience Function

↓

Pitch Contour Creation

↓

Melody Selection

Fraunhofer IDMT

Pitch Detection
Traditional Methods (Melodia)

- Melodia plugin available for Sonic Visualiser

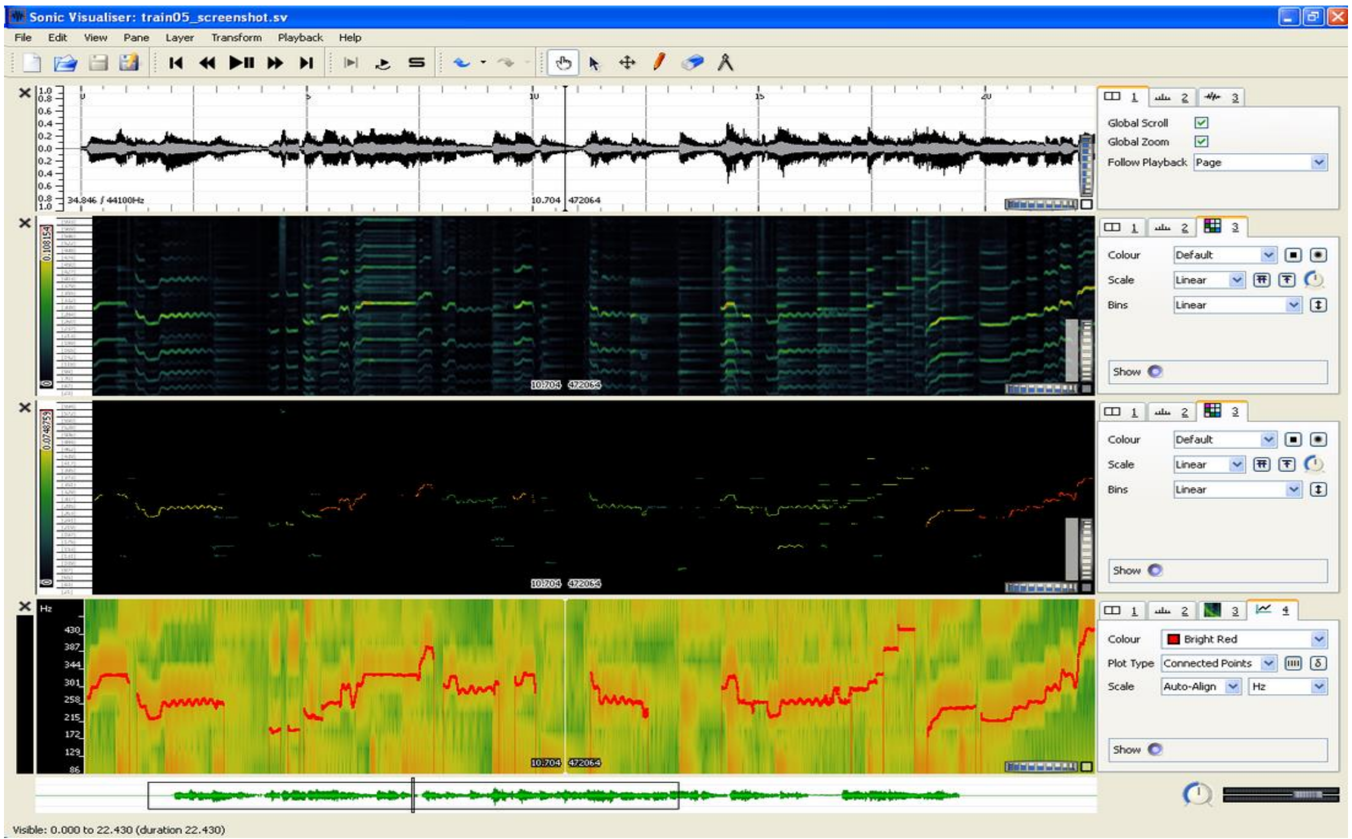

Fig. 6

Pitch Detection
Novel Methods

- CREPE (Convolutional Representation for Pitch Estimation) [Kim et al., 2018]
  - Monophonic pitch tracker
  - End-to-end modeling
    - Audio samples → pitch likelihoods
    - 20 cent resolution (5 pitch bins per semitones)

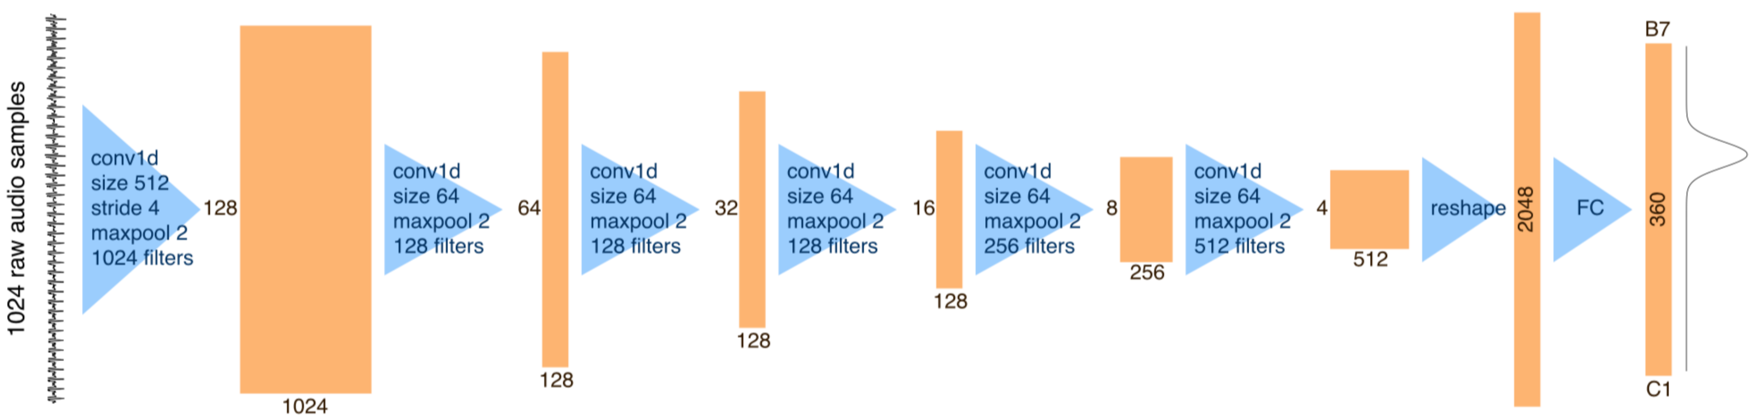

Fig. 7

Pitch Detection
Novel Methods

- Auto-encoder structure (U-Net) [Hsieh et al., 2019]
  - Mapping: multiple time-frequency representations (2D) → pitch saliency map (2D)
  - (Bottleneck) embedding encodes pitch voicing (melody activity)

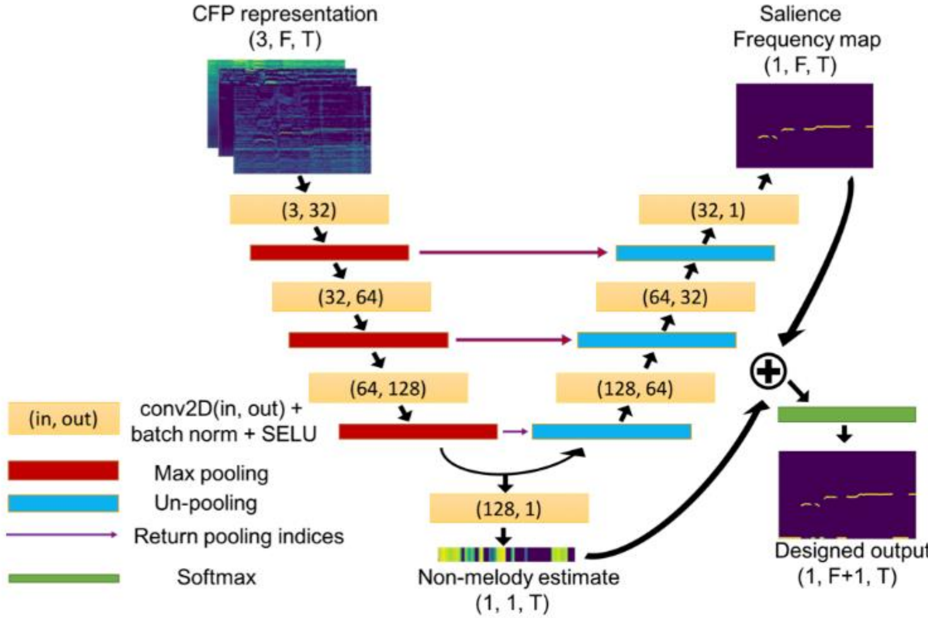

Fig. 8

Instrument Recognition
Introduction

- Music ensembles include multiple instruments
  - Sound production (string / wind / brass / drum instruments)
  - Unique timbre
- Overlapping sound sources (solo recording vs. orchestra)
  - Unison (same pitch)
  - Harmonic intervals (overtone overlap)
  - Rhythmic interconnection (note attacks overlap)
- Classification on different taxonomy levels
  - Woodwind instruments → saxophone → tenor saxophone

Instrument Recognition
Tasks

- Sorted by increasing complexity/difficulty
  - Instrument recognition of isolated note recordings
  - Instrument recognition on isolated instrument tracks
  - Predominant instrument recognition in ensemble recordings
  - Polyphonic instrument recognition (classify all instruments)

Instrument Recognition
Traditional Methods

- Multiple categories of audio features [Grasis et al., 2014]
    - Frame-level
        - e.g., spectral flux & flatness
    - Overtone-level
        - e.g., modulation rate & frequency
    - Note-event level
        - e.g., magnitude ratios of overtones

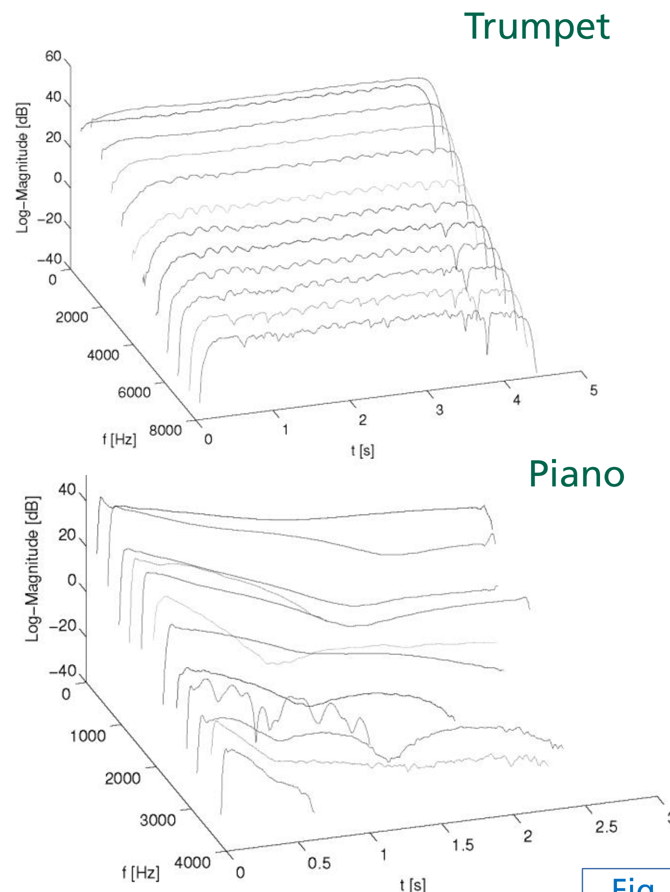

Fig. 9

Instrument Recognition
Novel Methods

- Mel spectrogram + CNN model [Han et al., 2017]

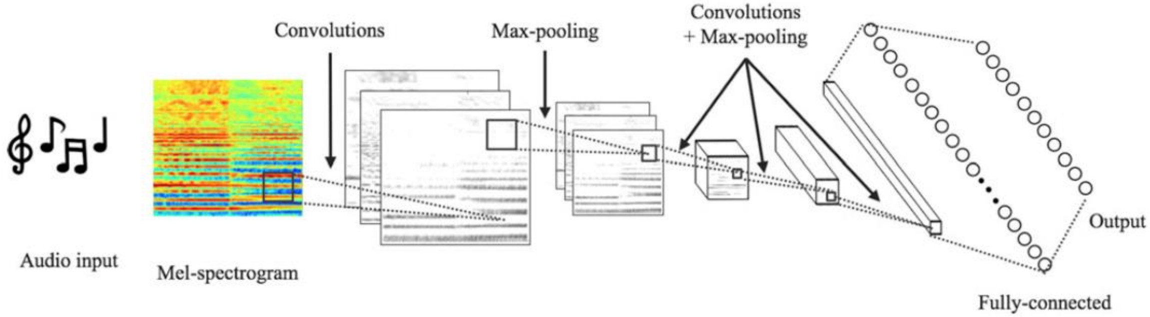

Fig. 10

- Improved separability of instrument classes in feature spaces

Deeper layers

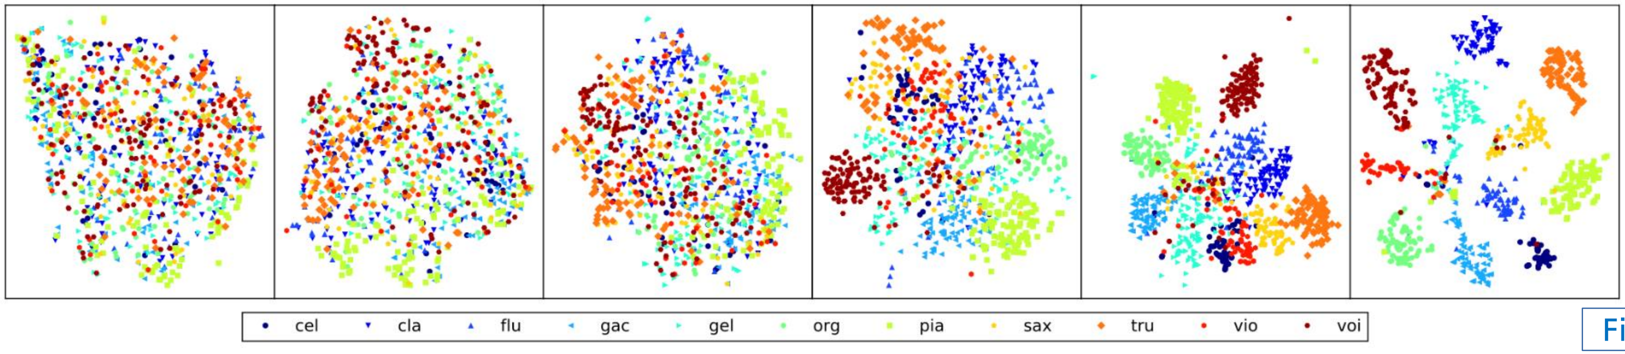

Fig. 11

Instrument Recognition
Novel Methods

- Pitch-Informed Frame-level Instrument Recognition [Hung & Yang, 2018]

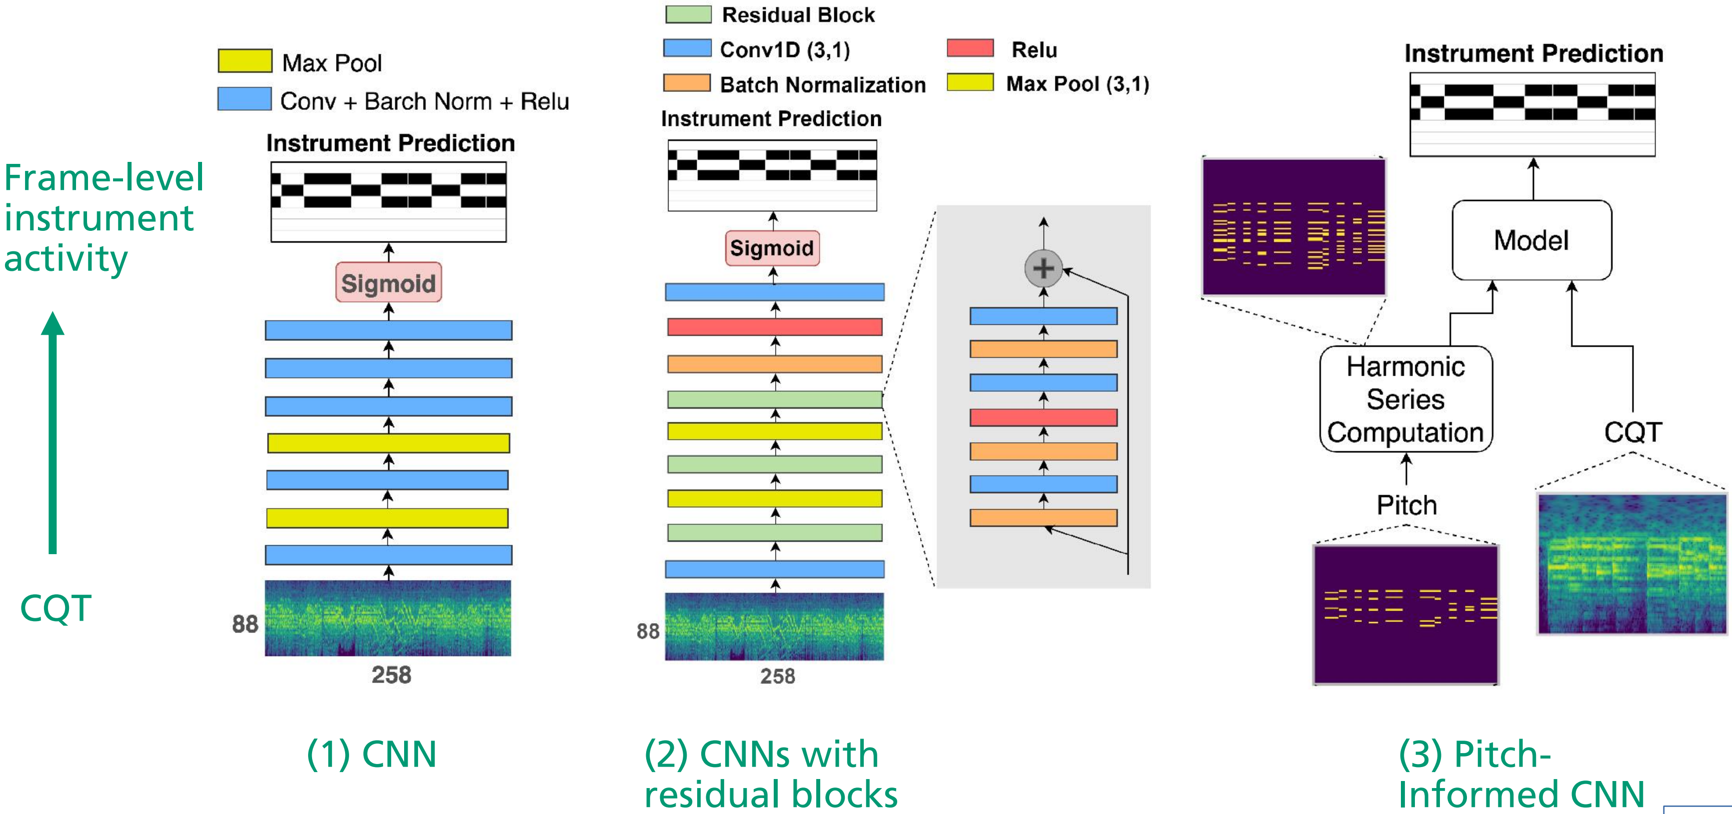

Fig. 12

Source Separation
Introduction

- Music recordings
  - Mixtures of different musical instruments (sources) playing simultaneously
- Sound Separation
  - Reverse engineering the audio mixing process
  - Output: 1 stem per instrument

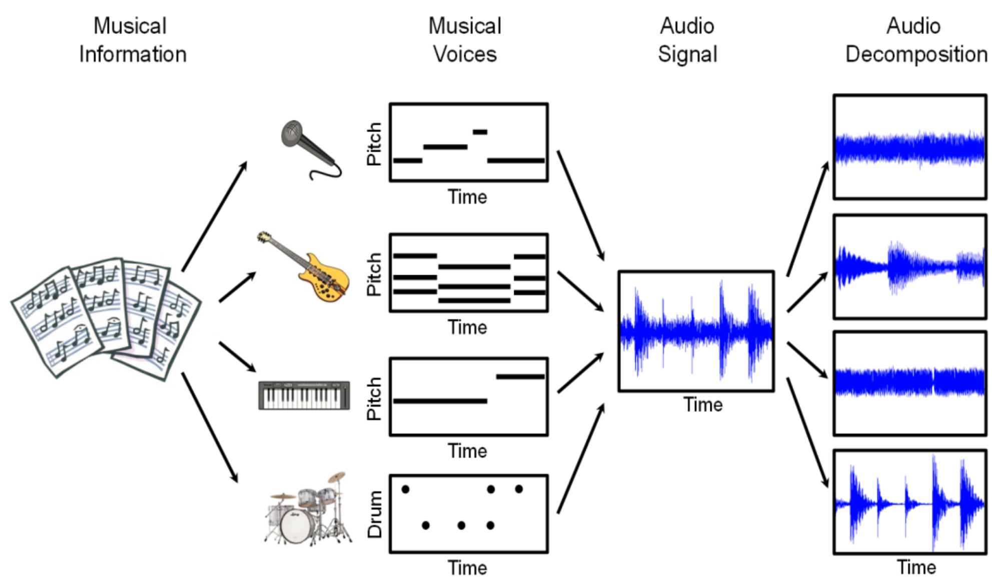

Fig. 18

Fraunhofer IDMT

Source Separation
Introduction

- Audio mix is influenced by
  - Instrument characteristics (timbre, note decay, …)
  - Musical performance (timing, dynamics, playing techniques, …)
  - Recording chain (microphones, room acoustics)
  - Post-processing (effects, mastering, DAW mix)

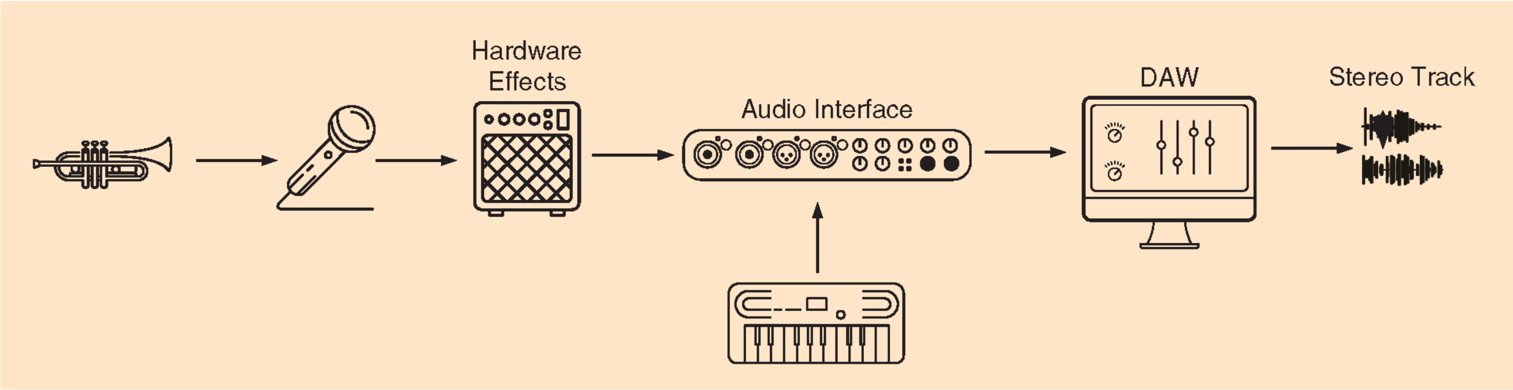

Fig. 13

Source Separation
Application Scenarios

- Audio remixing
- Audio upmixing
    - Mono → stereo
    - Stereo → 5.1
- Music Analysis
    - Transcription, beat tracking, harmony analysis etc.
- Music Education
    - Solo / Backing track generation

Source Separation
Tasks

- Harmonic/percussive separation
  - H → stable harmonic components (fundamental frequency, overtones)
  - P → transient components (drum sounds, note attacks)
- Solo/accompaniment separation
  - S → predominant melody instrument
  - A → accompanying instruments
- Singing voice separation
  - S → singing voice (male / female)
  - A → band
- Separation of all sources

Fraunhofer IDMT

Source Separation
Traditional Approaches

- Harmonic/percussive (H/P) separation
  - Different spectral characteristics of harmonic and percussive signals

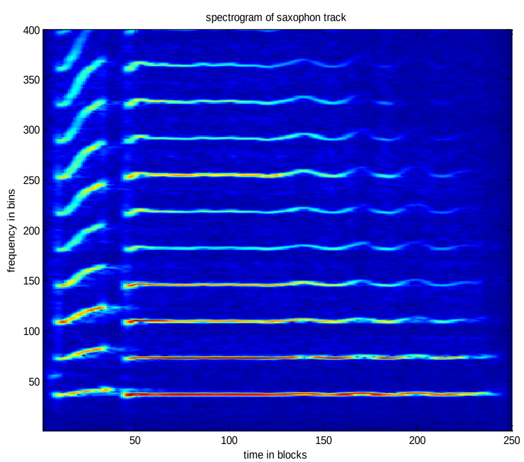

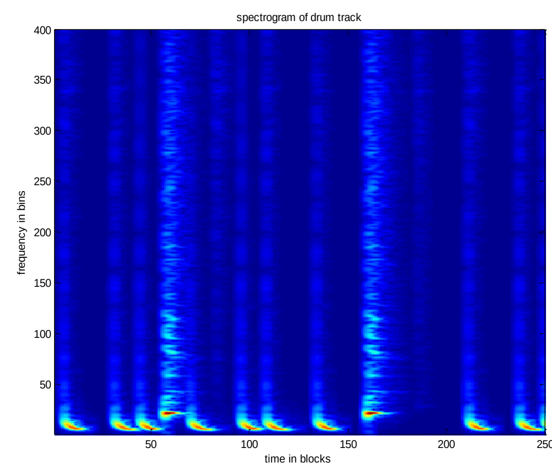

- Time-continuous (horizontal)
- Localized in frequency

- Wide-band (vertical)
- Localized in time

Source Separation
Traditional Approaches

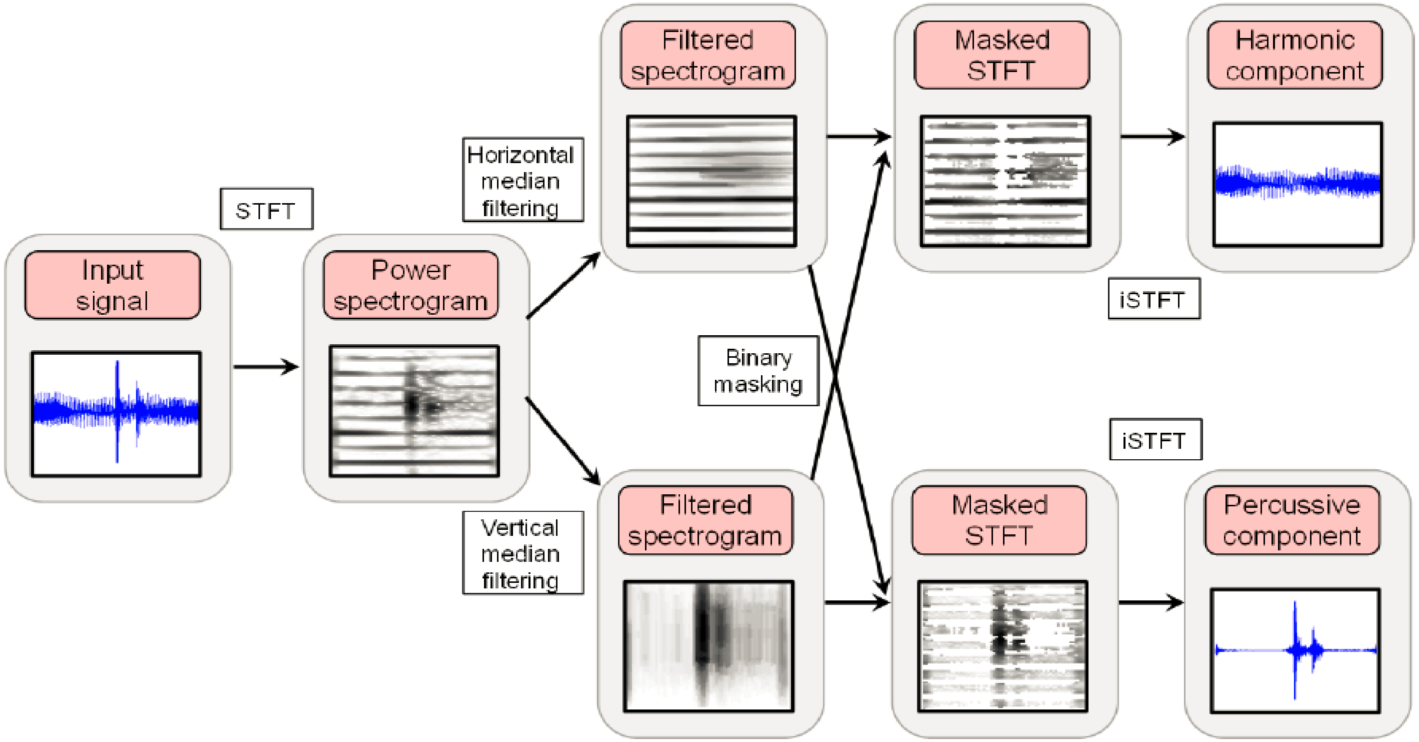

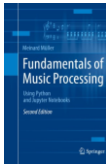

FMP Notebooks

Fig. 14

Source Separation
Traditional Approaches

- Phase-based H/P separation
  - Harmonic sources → phase change values are predictable
  - Percussive sources → unpredictable phase (noise-like characteristics)
  - Instantaneous Frequency Distribution (IFD)
    - How does phase change over time?

$$\Phi(k,n) = \frac{1}{2\pi}\frac{d\phi(k,n)}{dn}$$

Unwrapped phase

Instantaneous Frequency

k – frequency bin

n – time frame

Source Separation
Traditional Approaches

- Phase-based H/P separation
  - Harmonic mask → phase change within range / predictable?

$$H(k,n) = \begin{cases} 1 & \text{if } \Delta_{k_{Low}} < \Phi(k,n) < \Delta_{k_{High}} \\ 0 & \text{otherwise} \end{cases}$$

  - Percussive mask

$$P(k,n) = 1 - H(k,n)$$

Source Separation
Traditional Approaches

- Non-Negative Matrix Factorization (NMF)
  - Factorize spectrogram *V* into set of components: $V \approx WH$

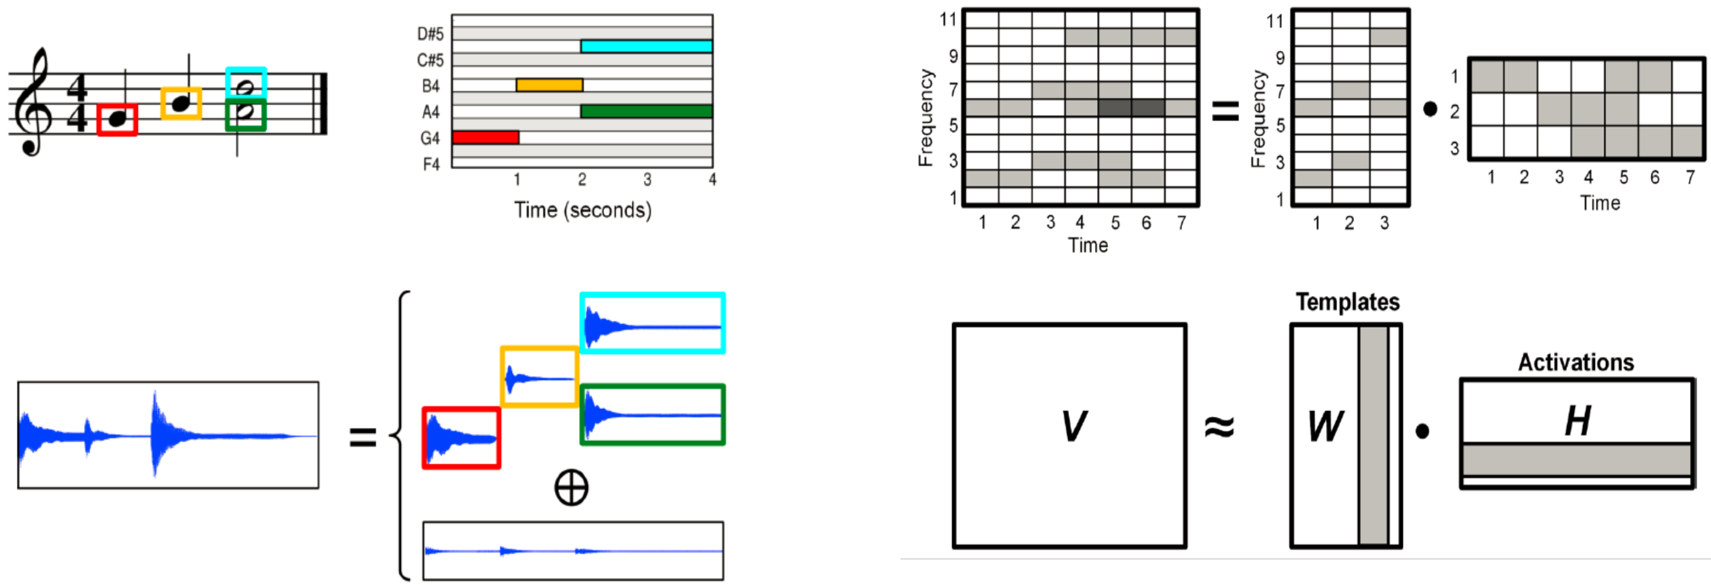

Source Separation
Traditional Approaches

- Non-Negative Matrix Factorization (NMF)
  - Algorithm: $V \approx WH$
    - Randomly initialize W & H
    - Use update rules to alternately update W & H
      - Minimize cost function
    - Cost function examples
      - Euclidean distance

$$||A - B||^2 = \sum_{ij} (A_{ij} - B_{ij})^2$$

      - Kullback-Leibler divergence

$$D(A||B) = \sum_{ij} \left(A_{ij} \log \frac{A_{ij}}{B_{ij}} - A_{ij} + B_{ij} \right)$$

Fraunhofer IDMT

Source Separation
Novel Approaches

- U-Net based [Jannson et al., 2017]
  - Input → magnitude spectrogram (mix)
  - Output → 2 soft masks (voice / others)
- Issue
  - Decoupling of phase and magnitude

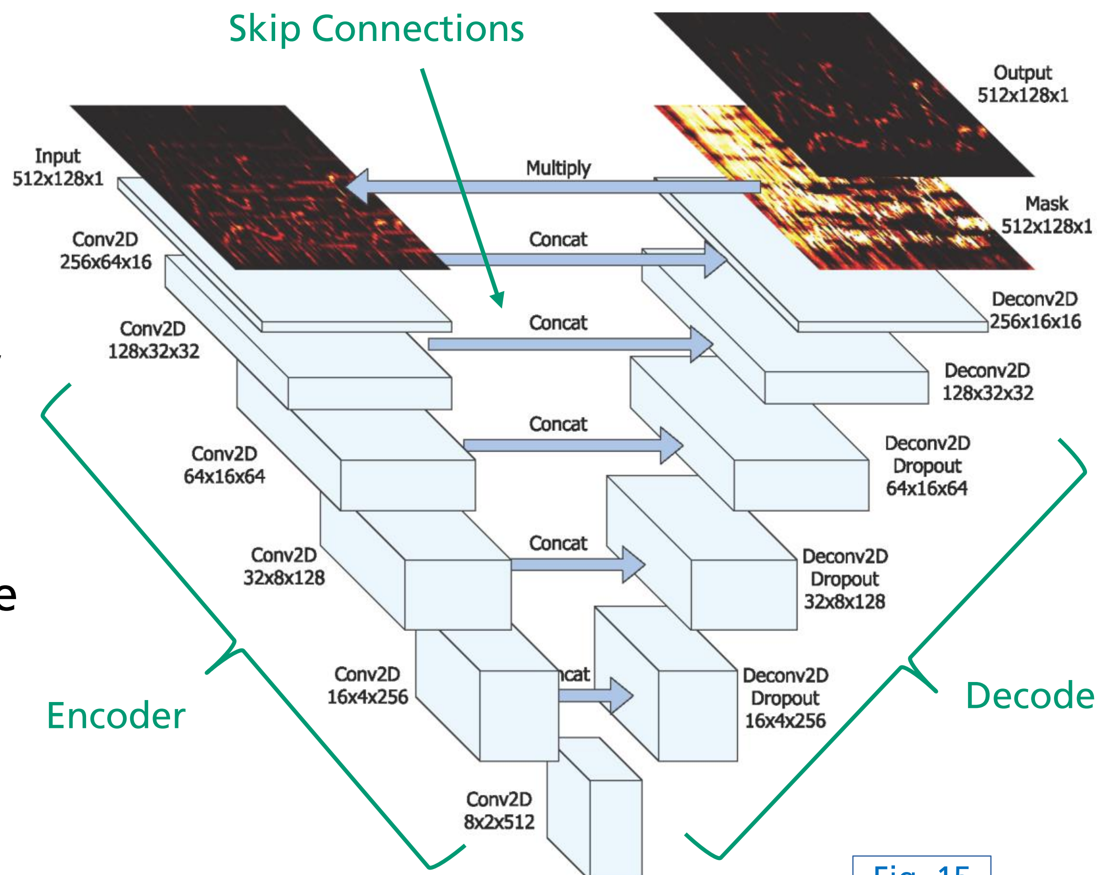

Fig. 15

Source Separation
Novel Approaches

- Spleeter [Hennequin et al., 2020]
  - Open-source version for MIR research
  - 3 pre-trained models
    - 2 stems (vocals and accompaniments)
    - 4 stems (vocals, drums, bass, and other)
    - 5 stems (vocals, drums, bass, piano and other)

Spleeter Demo

Source Separation
Novel Approaches

- Conv-TasNet [Luo & Mesgarani, 2019]
  - Time-domain speech separation network (end-to-end)
  - Encoder → optimized representation for speaker separation
  - Seperation → masks (weighting functions)
  - Decoder → invert to waveforms
  - Temporal convolutional networks (TCN)
    - Stack of 1-D dilated convolutional blocks
    - Large receptive field → model long-term dependencies

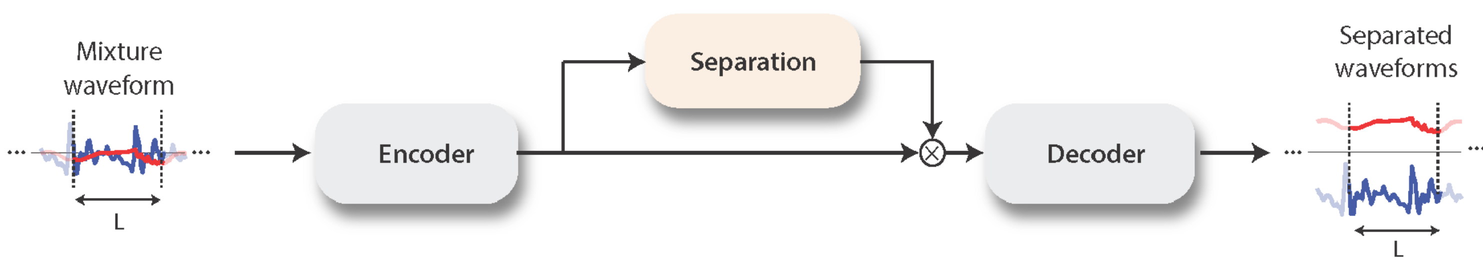

Fig. 16

Source Separation
Novel Approaches

- Conv-TasNet [Luo & Mesgarani, 2019]

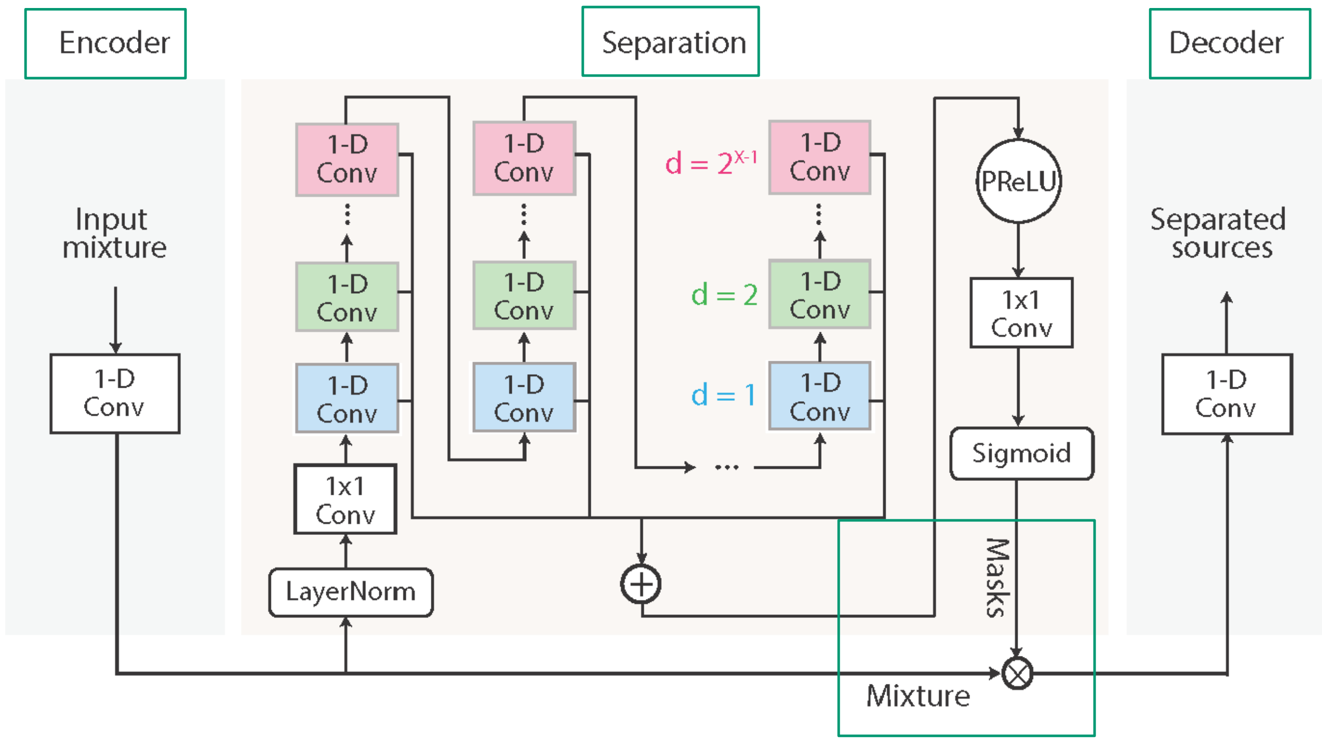

Fig. 17

Summary

- Case Studies
    - Pitch Detection
    - Instrument Recognition
    - Source Separation

References

Cano, E., Fitzgerald, D., Liutkus, A., Plumbley, M. D., & Stoter, F. R. (2019). Musical Source Separation: An Introduction. *IEEE Signal Processing Magazine*, *36*(1), 31–40.

Grasis, M., Abeßer, J., Dittmar, C., & Lukashevich, H. (2014). A Multiple-Expert Framework for Instrument Recognition. *Lecture Notes in Computer Science 8905*, 619–634.

Han, Y., Kim, J., & Lee, K. (2017). Deep Convolutional Neural Networks for Predominant Instrument Recognition in Polyphonic Music. *IEEE/ACM Transactions on Audio Speech and Language Processing*, *25*(1), 208–221.

Hennequin, R., Khlif, A., Voituret, F., & Moussallam, M. (2020). Spleeter: a fast and efficient music source separation tool with pre-trained models. *Journal of Open Source Software*, *5*(50), 2154.

Hsieh, T. H., Su, L., & Yang, Y. H. (2019). A Streamlined Encoder/Decoder Architecture for Melody Extraction. *Proceedings of the IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 156–160. Brighton, UK.

Hung, Y.-N., & Yang, Y.-H. (2018). Frame-Level Instrument Recognition by Timbre and Pitch. *Proceedings of the International Society for Music Information Retrieval Conference (ISMIR)*, 135–142. Paris, France.

Jansson, A., Humphrey, E., Montecchio, N., Bittner, R., Kumar, A., & Weyde, T. (2017). Singing Voice Separation with Deep U-Net Convolutional Networks. *Proceedings of the International Society for Music Information Retrieval Conference (ISMIR)*, 745–751. Suzhou, China.

References

Kim, J. W., Salamon, J., Li, P., & Bello, J. P. (2018). Crepe: A Convolutional Representation for Pitch Estimation. *Proceedings of the IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 161–165. New Orleans, USA.

Luo, Y., & Mesgarani, N. (2019). Conv-TasNet: Surpassing Ideal Time-Frequency Magnitude Masking for Speech Separation. *IEEE/ACM Transactions on Audio, Speech, and Language Processing*, *27*(8), 1256–1266.

Müller, M. (2021). *Fundamentals of Music Processing - Using Python and Jupyter Notebooks* (2nd ed.). Springer.

Salamon, J., & Gomez, E. (2012). Melody extraction from polyphonic music signals using pitch contour characteristics. *IEEE Transactions on Audio, Speech and Language Processing*, *20*(6), 1759–1770.

Images

Fig. 1: [Müller, 2021], p. 449, Fig. 8.15(b)

Fig. 2: http://www.guitaradventures.com/wp-content/uploads/Tuning-your-guitar.jpg

Fig. 3: https://cdn2.whatoplay.com/screenshots/2631slide-4.jpg

Fig. 4: https://cdn.androidcommunity.com/wp-content/uploads/2010/11/500x_angrybirdsdarwin.jpg

Fig. 5: [Müller, 2021], p. 449, Fig. 8.15(a)

Fig. 6: Sonic Visualiser: http://www.sonicvisualiser.org/ , Melodia plugin: http://mtg.upf.edu/technologies/melodia

Fig. 7: [Kim et al., 2018], p. 2, Fig. 1

Fig. 8: [Hsieh et al., 2019], p. 2, Fig. 2

Fig. 9: [Grasis et al., 2014], p. 6, Fig. 3

Fig. 10: [Han et al., 2017], p. 3, Fig. 1

Fig. 11: [Han et al., 2017], p. 9, Fig. 6

Fig. 12: [Hung & Yang, 2018], p. 4, Fig. 1

Fig. 13: [Cano et al., 2019], p. 3, Fig. 3

Fig. 14: [Müller, 2021], p. 425, Fig. 8.3

Images

Fig. 15: [Jansson, 2017], p. 3, Fig. 1

Fig. 16: [Luo & Mesgarani, 2019], p. 3, Fig. 1(A)

Fig. 17: [Luo & Mesgarani, 2019], p. 3, Fig. 1(B)

Fig. 18: [Müller, 2021], p. 422, Fig. 8.1

Sounds

AUD-1: Aislinn – Capclear (2013), https://freemusicarchive.org/music/Aislinn/Aislinn/10_-_Aislinn_-_Capclear

AUD-2: Aislinn – Fourteen Days (2013), https://freemusicarchive.org/music/Aislinn/Aislinn/11_-_Aislinn_-_Fourteen_days

AUD-3: Anonymous Choir – Amicus Meus (2009), https://freemusicarchive.org/music/Anonymous_Choir/Toms_Luis_de_Victorias_Amicus_Meus/Amicus_Meus

Fraunhofer IDMT

Thank you!

- Any questions?

Dr.-Ing. Jakob Abeßer

Fraunhofer IDMT

Jakob.abesser@idmt.fraunhofer.de

https://machinelistening.github.io

Fraunhofer IDMT